

# Fast Bilateral-Space Stereo for Synthetic Defocus

Jonathan T. Barron  
Andrew Adams  
YiChang Shih  
Carlos Hernández

Google

# Shallow Depth of Field



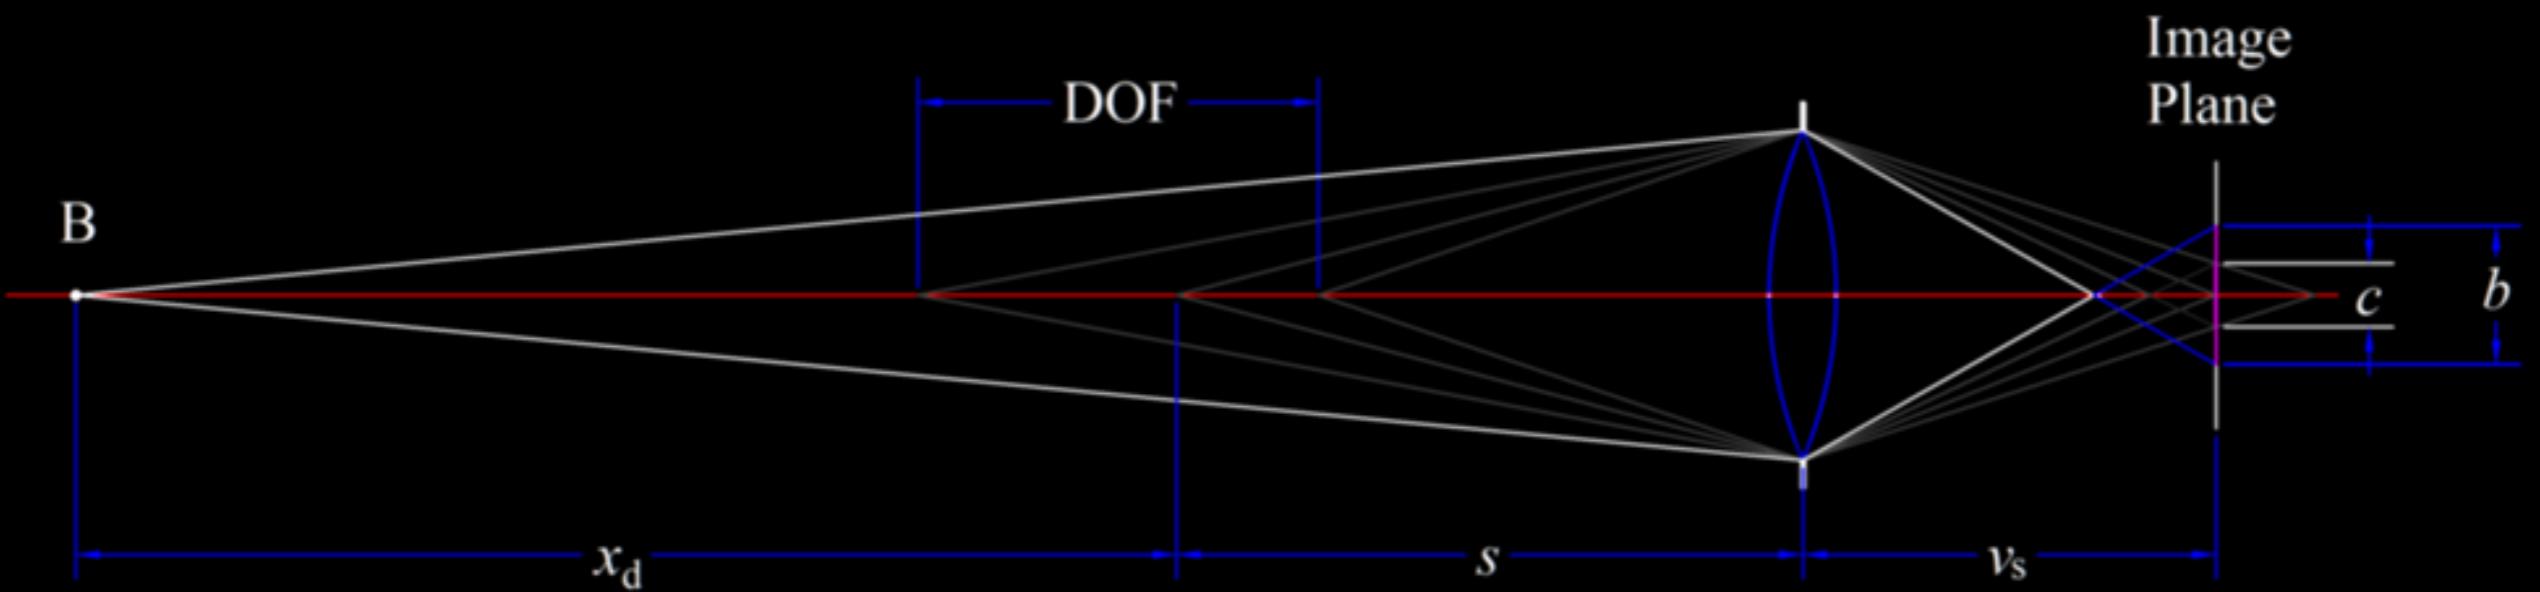
Photo by Bryan Maltais

# Deep Depth of Field



Photo by Bryan Maltais

# Real Shallow Depth of Field



# Synthetic Shallow Depth of Field



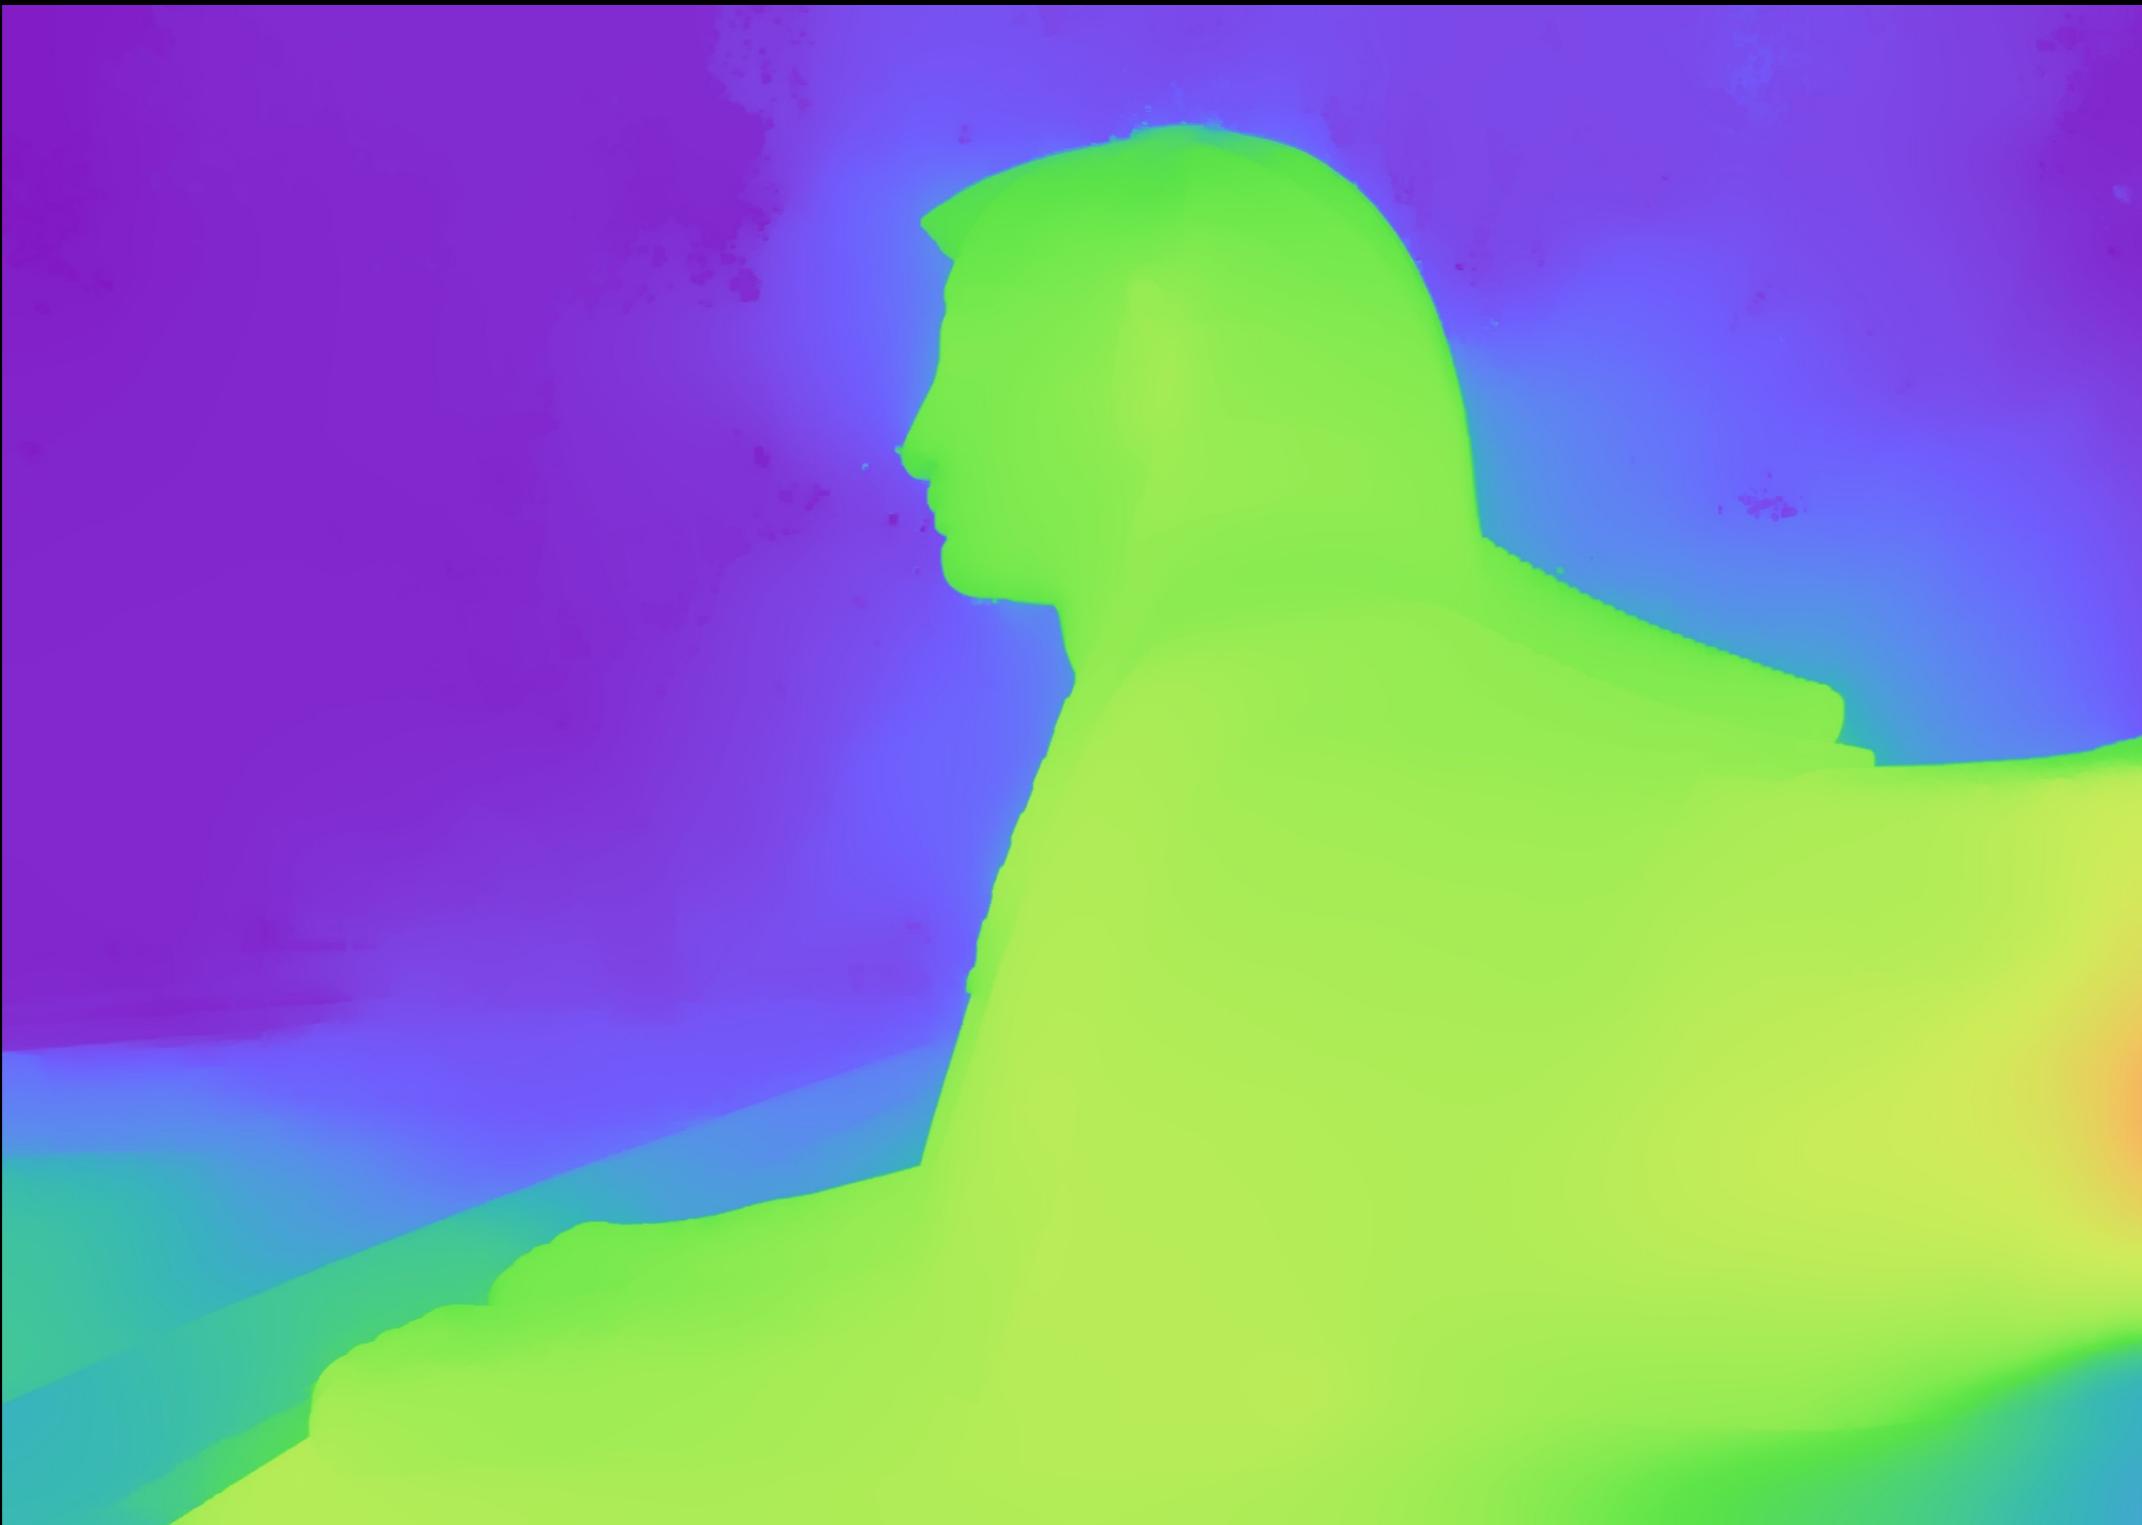
Stereo Pair

# Synthetic Shallow Depth of Field



Stereo Pair

# Synthetic Shallow Depth of Field



Depth Map

# Synthetic Shallow Depth of Field



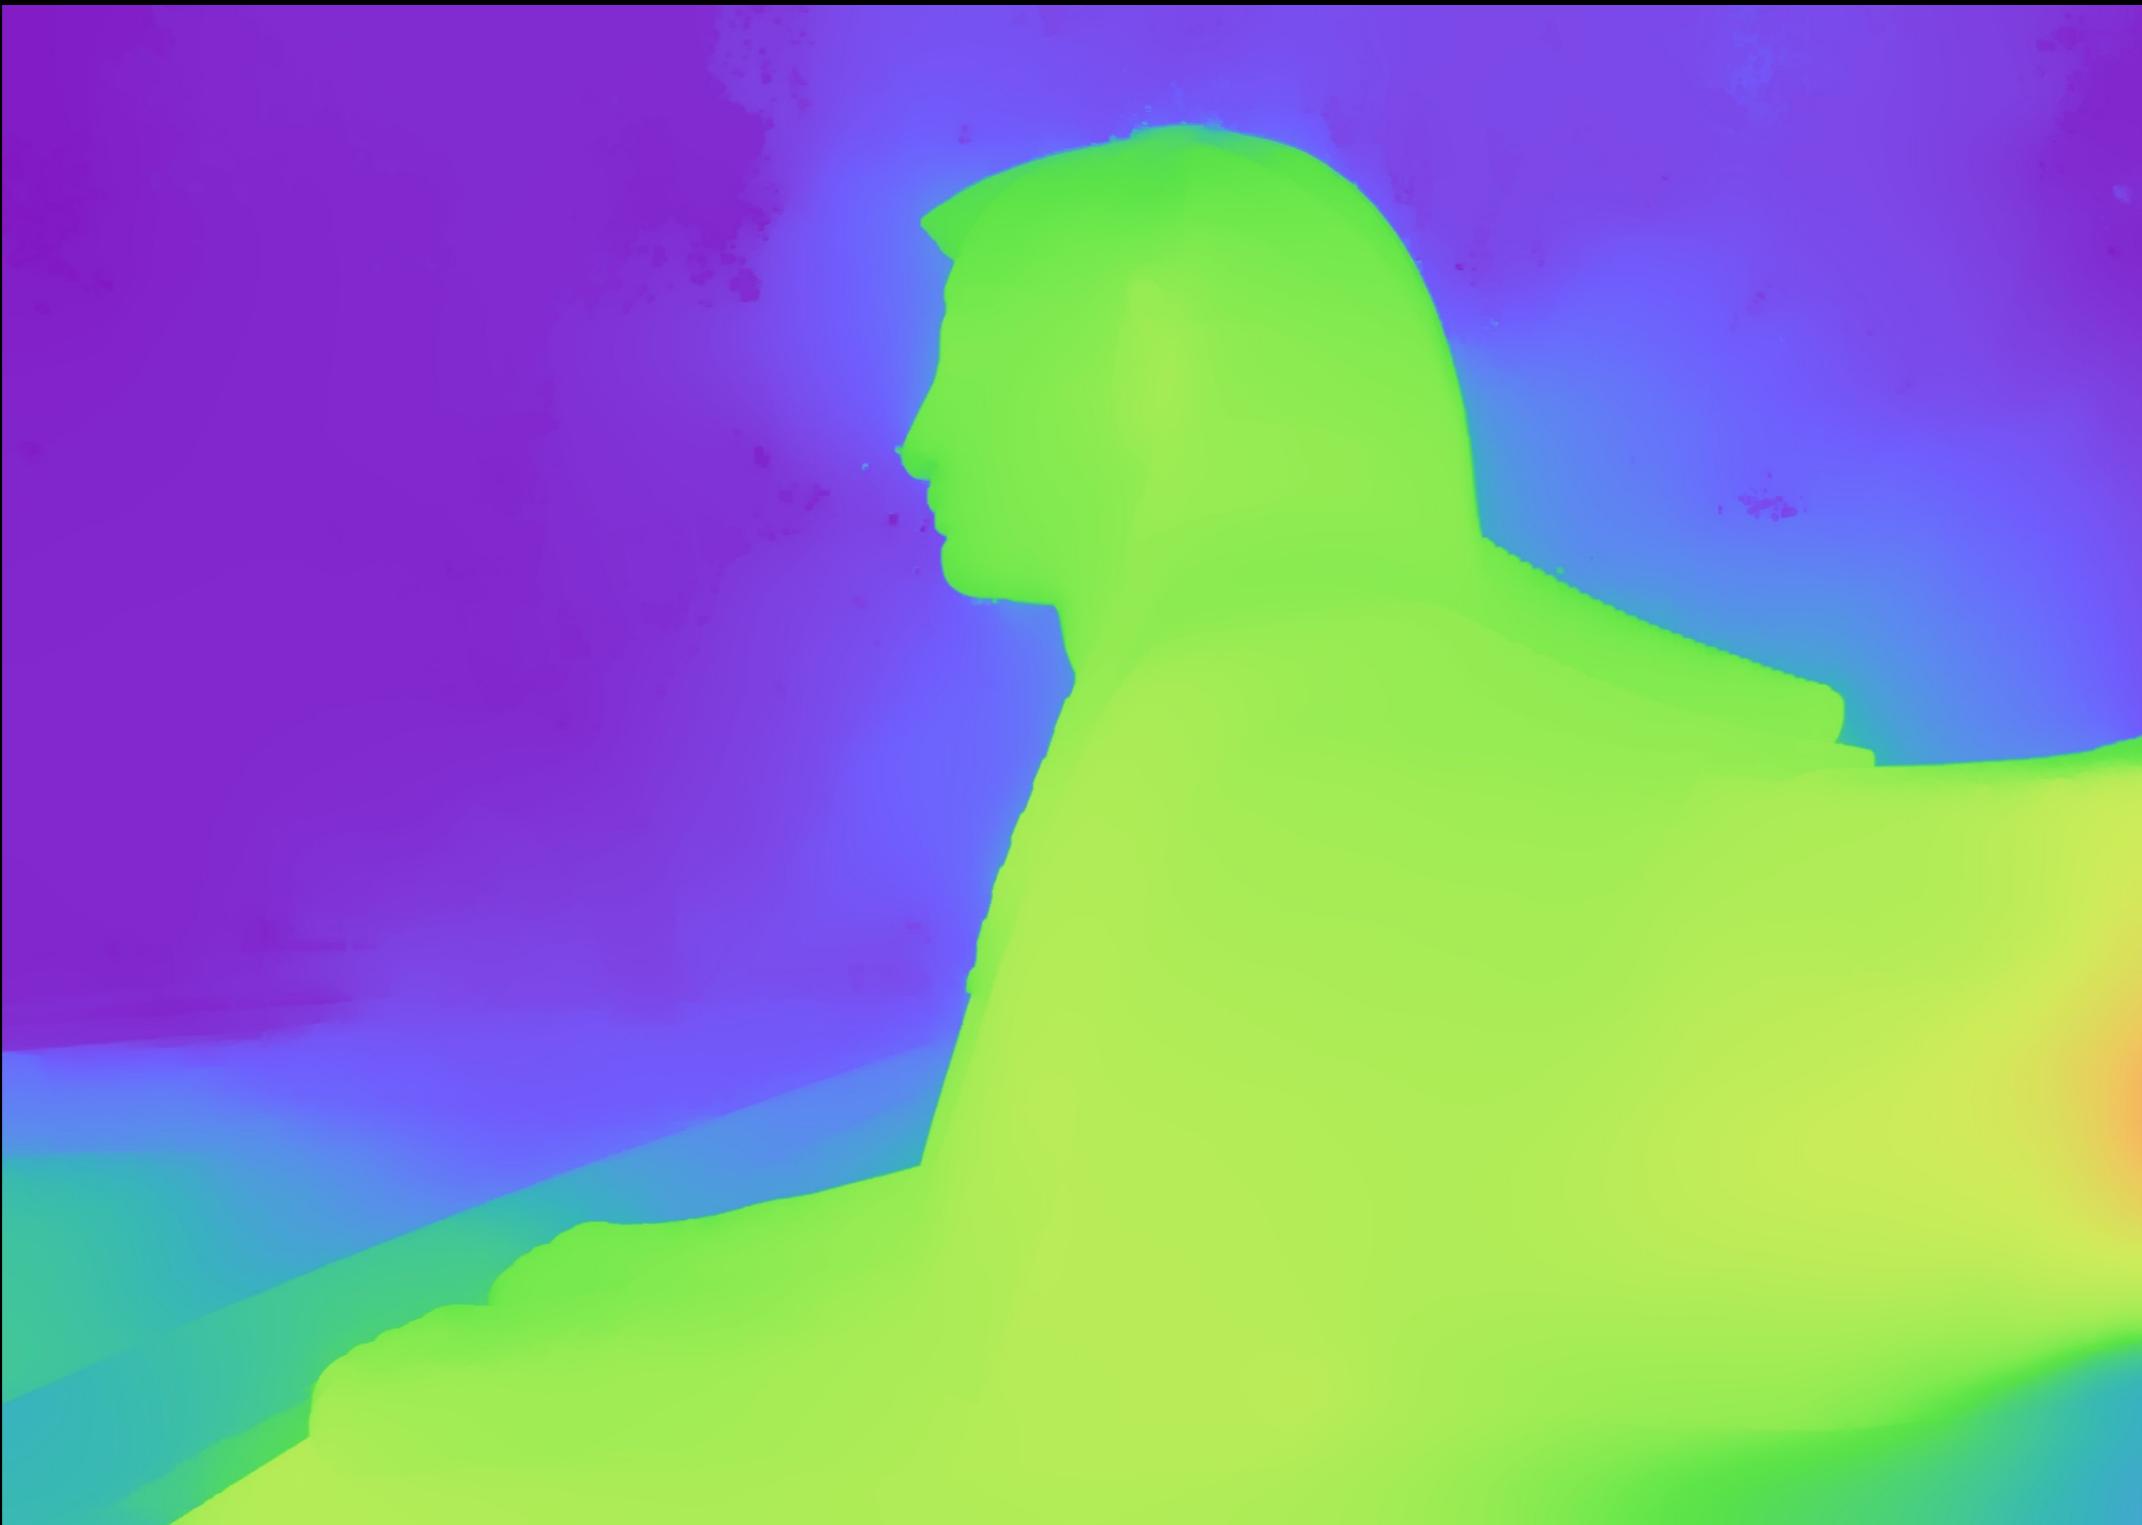
Rendered Image

# Synthetic Shallow Depth of Field



Input Image

# Synthetic Shallow Depth of Field



Depth Map

# “Isn’t stereo solved?”

**Speed:** Must be fast, even on a mobile device, and even on huge images

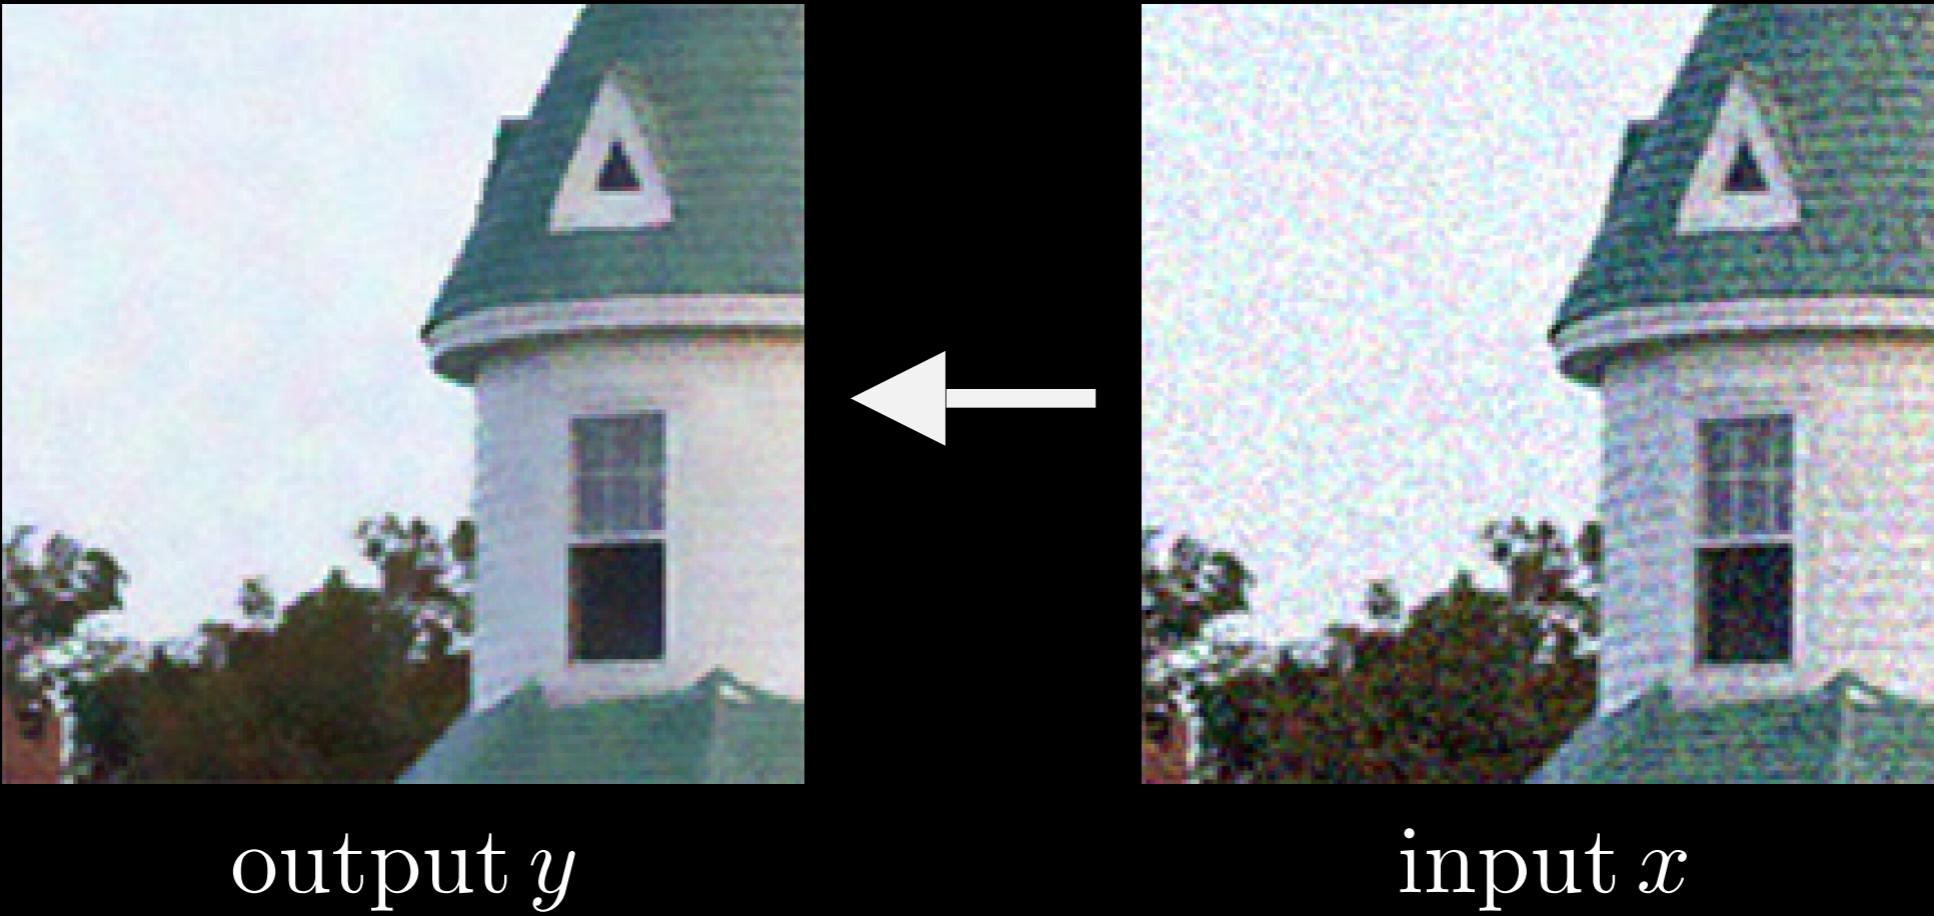
Our algorithm is **10-100x faster** than the state-of-the-art

**Quality:** State-of-the-art stereo algorithms only roughly follow edges in the input image

Our algorithm **cannot** produce depth maps which are mis-aligned to the input image

# Review:

# Bilateral Filtering



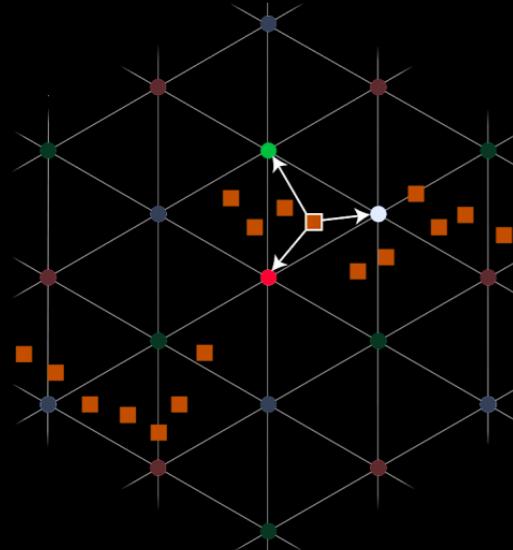
Blur within objects, not across edges

$$y = Ax$$

$$A_{i,j} = \exp \left( -\frac{\|[x_i, y_i] - [x_j, y_j]\|^2}{\sigma_{xy}^2} - \frac{\|[r_i, g_i, b_i] - [r_j, g_j, b_j]\|^2}{\sigma_{rgb}^2} \right)$$

# Review: Fast Bilateral Filtering

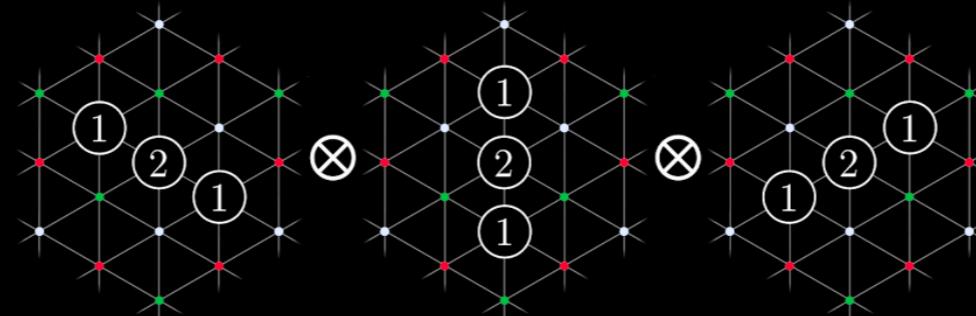
Input



Splat

resample from pixels  
into “bilateral-space”

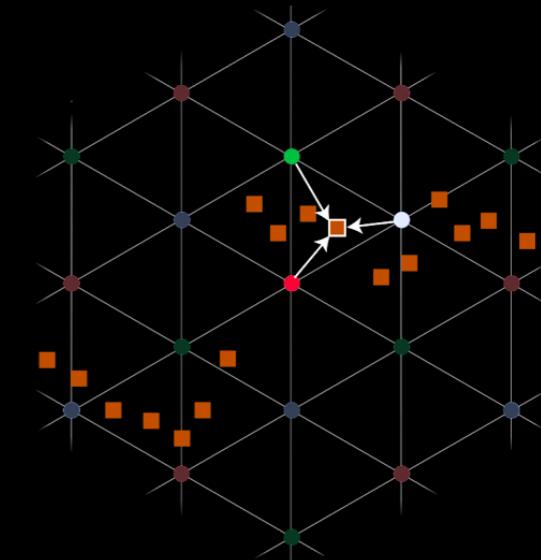
local, cheap filtering in  
“bilateral-space”



Blur

apply a series of blurs  
in bilateral-space

Output



Slice

resample back  
into pixel-space

=

Non-local, edge-aware  
filtering in pixel-space

# Past Work

To quickly apply a bilateral filter:

- 1) splat the signal
- 2) blur the splatted signal
- 3) slice out the blurred signal

# Our Insight

To quickly solve an optimization problem based on bilateral filtering:

- 1) splat the **problem**
- 2) solve the splatted problem
- 3) slice out the solution

# Bilateral-Space = Reparametrization



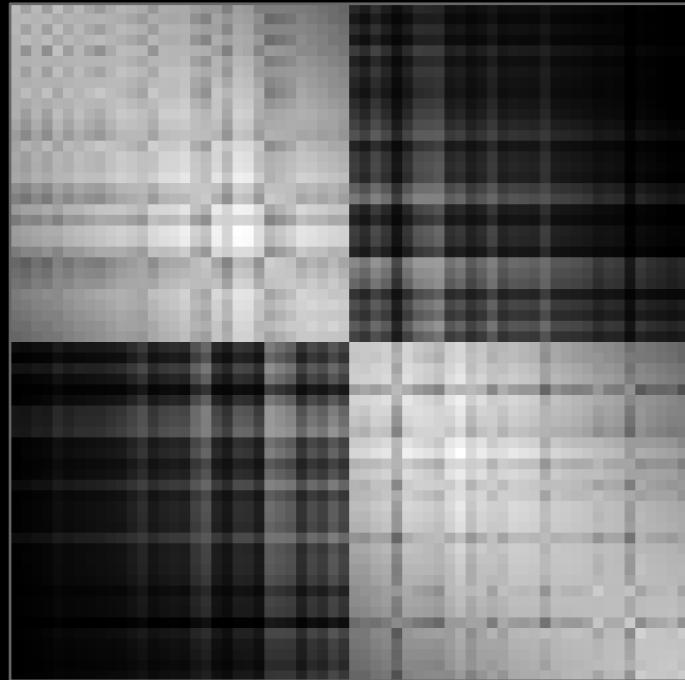
Instead of solving in “pixel space”  
(each color is a parameter)

solve in “bilateral-space”  
(each color is a parameter)

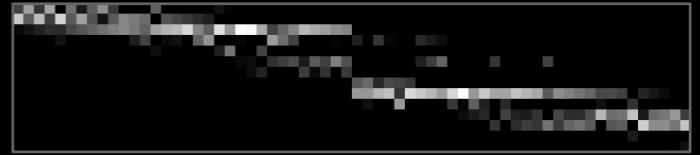
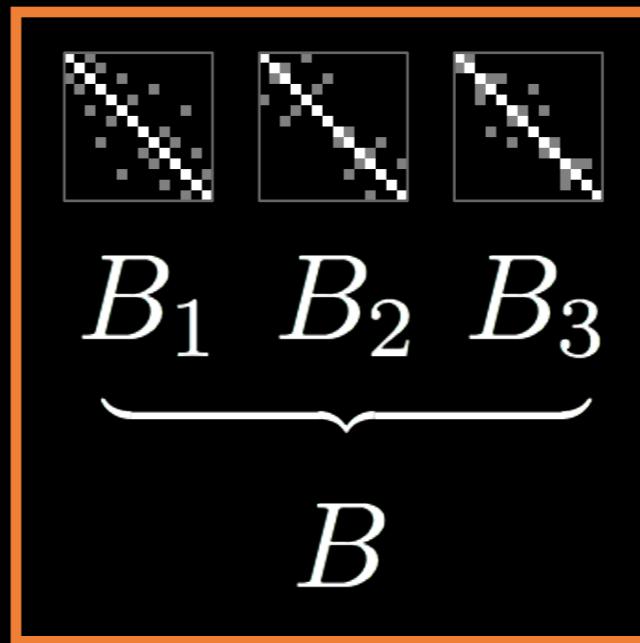
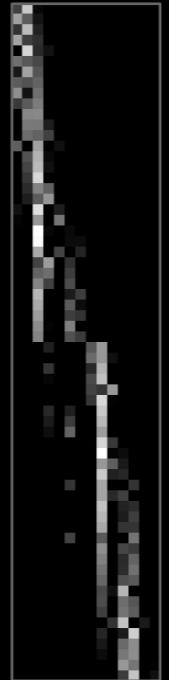
- + Much smaller space
- + Long-range affinities are cheap
- + Solution must be bilateral-smooth

# Bilateral-Space = Matrix Decomposition

$$A = S^T B S$$



=



$S$

$A$

$S^T$

Affinity Matrix  
(Huge, Dense)

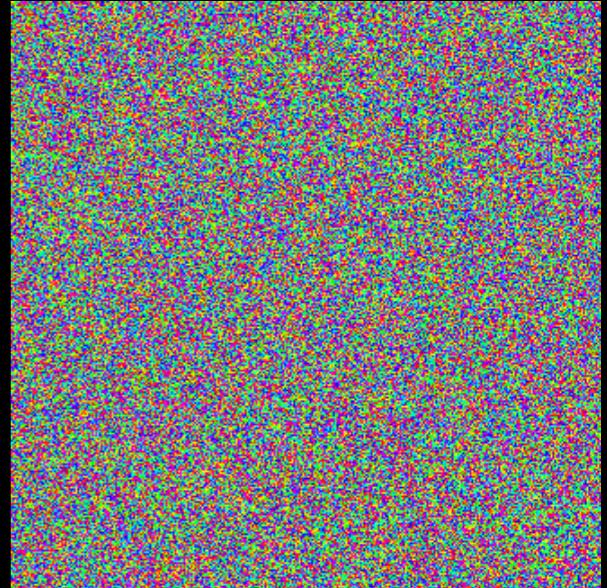
“Slice” Matrix  
(Skinny, Sparse)

“Blur” Matrices  
(Small, Sparse)

“Splat” Matrix  
(Skinny, Sparse)

# Pixel-space optimization

$$\underset{\mathbf{x}}{\text{minimize}} \quad \underbrace{\sum_i \sum_j A_{i,j} (x_i - x_j)^2}_{\text{The disparity should be bilateral-smooth}} + \lambda \sum_i f_i(x_i)$$



The disparity  
should be  
bilateral-smooth

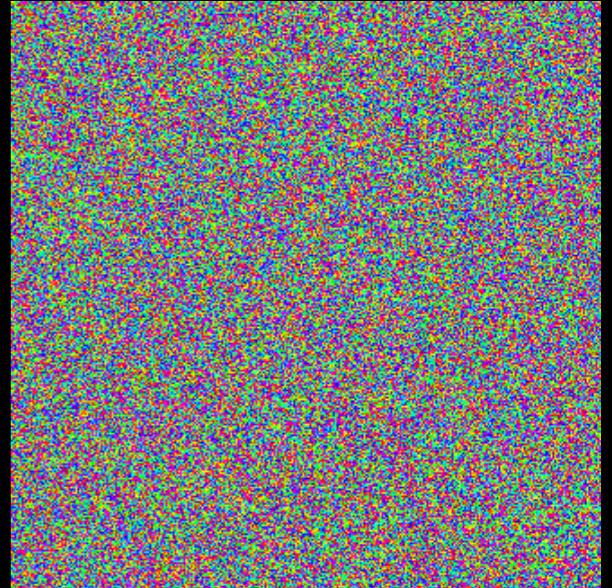
The warped right  
image should match  
the left image

- As many parameters as there are pixels (4-40 million)
- Loss function is expensive to evaluate and minimize

Slow at low resolutions, **intractable at high resolutions**

# Pixel-space optimization

$$\underset{\mathbf{x}}{\text{minimize}} \quad \sum_i \sum_j A_{i,j} (x_i - x_j)^2 + \lambda \sum_i f_i(x_i)$$

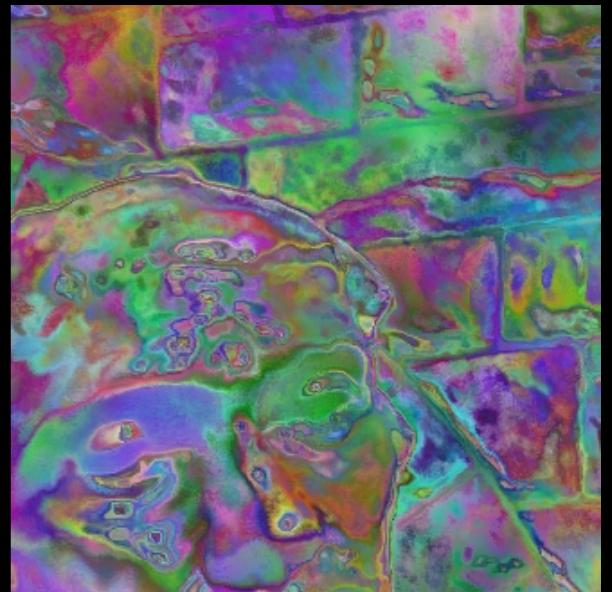


$\mathbf{m} := S\mathbf{1}$        $\forall x \quad g_j(x) = \sum_{(w,i) \in S_j} w f_i(x)$        $\mathbf{p} = S^T \mathbf{v}$   
 $\mathbf{n} := \vec{1}$        $f_i(p_i) = \max(0, p_i - u_i) + \max(0, l_i - p_i)$        $(C'_n \bar{B} C'_n) \mathbf{1} = (C_n \bar{B} C_n) \mathbf{1} = \mathbf{m}$   
While not converged       $g_j(v_j) = \sum_{y=0}^{v_j} \sum_{x=0}^y u''_j(x) + \sum_{y=v_j}^D \sum_{x=y}^{\infty} l''_j(x)$        $C'_n \approx C_n = \text{diag}(\mathbf{n})$   
 $\mathbf{n} := \sqrt{(\mathbf{n} \times \mathbf{m}) / (\bar{B} \mathbf{n})}$        $C'_s \approx C_s = \text{diag}(\mathbf{m})$



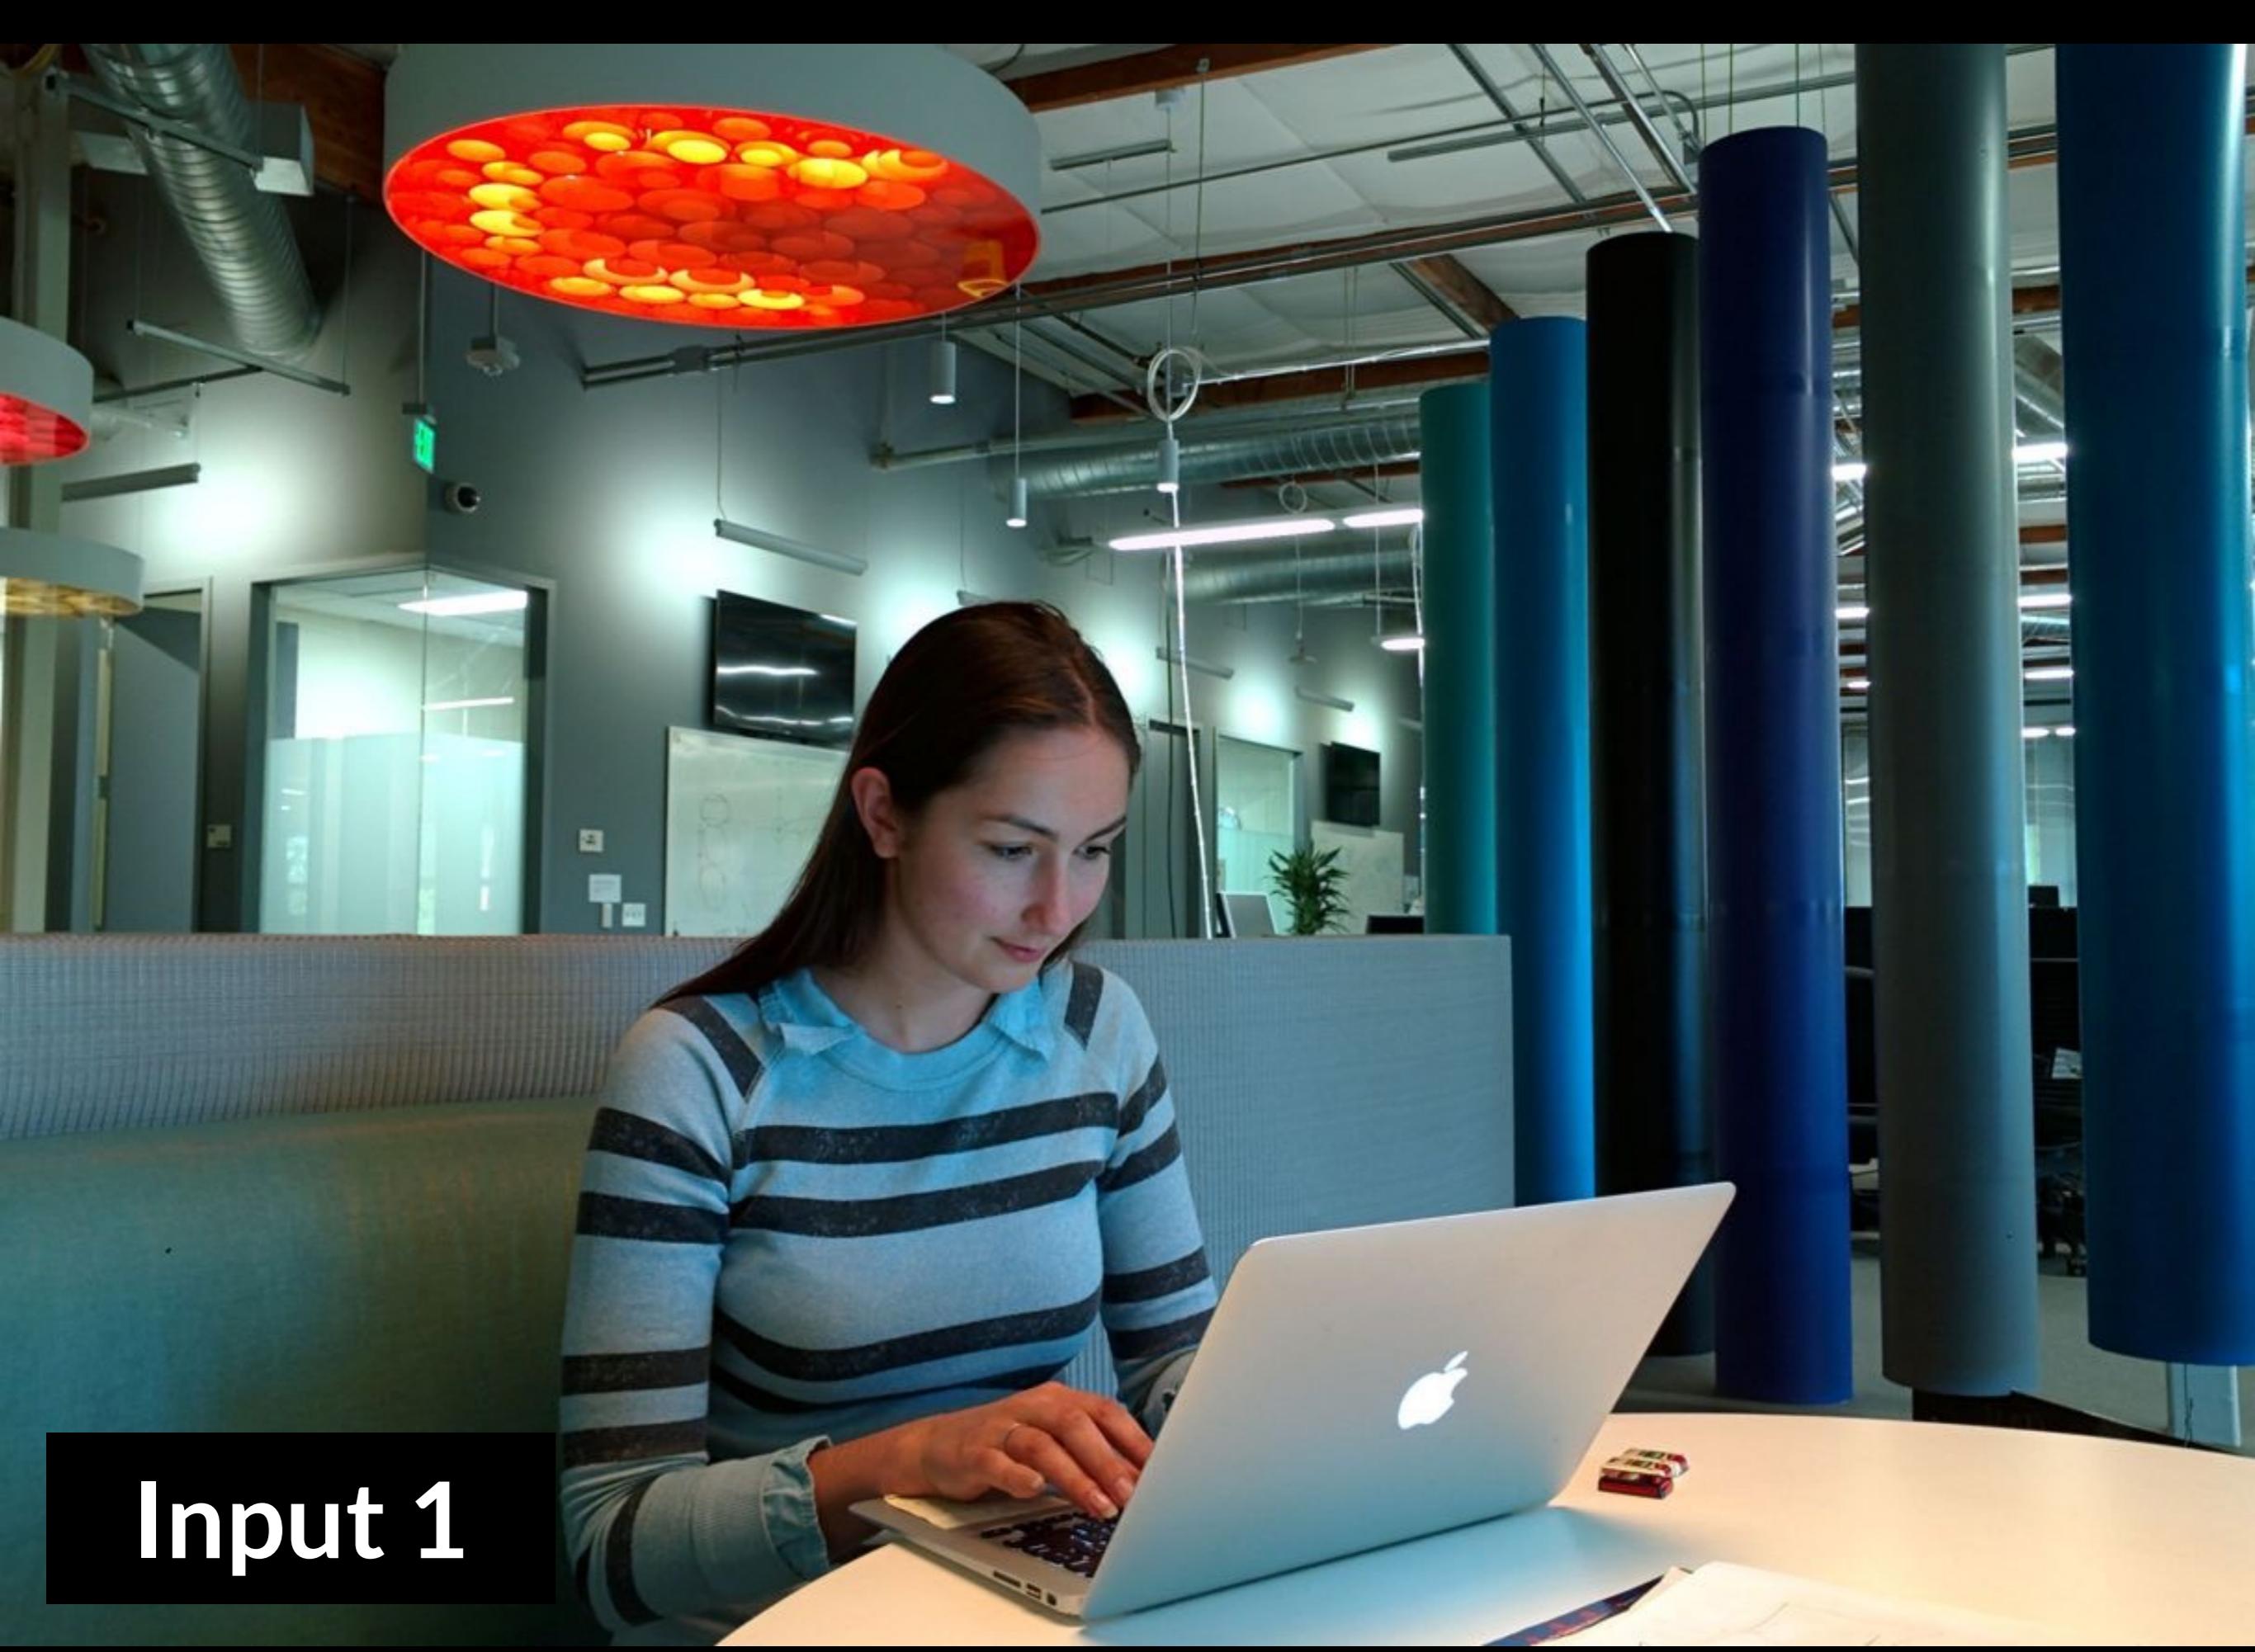
# Bilateral-space optimization

$$\underset{\mathbf{z}}{\text{minimize}} \quad \mathbf{z}^T (D - C B C) \mathbf{z} + \lambda \sum_j g_j(z_i)$$



- + Far fewer parameters (50-100x)
- + Loss function is very cheap to evaluate and minimize

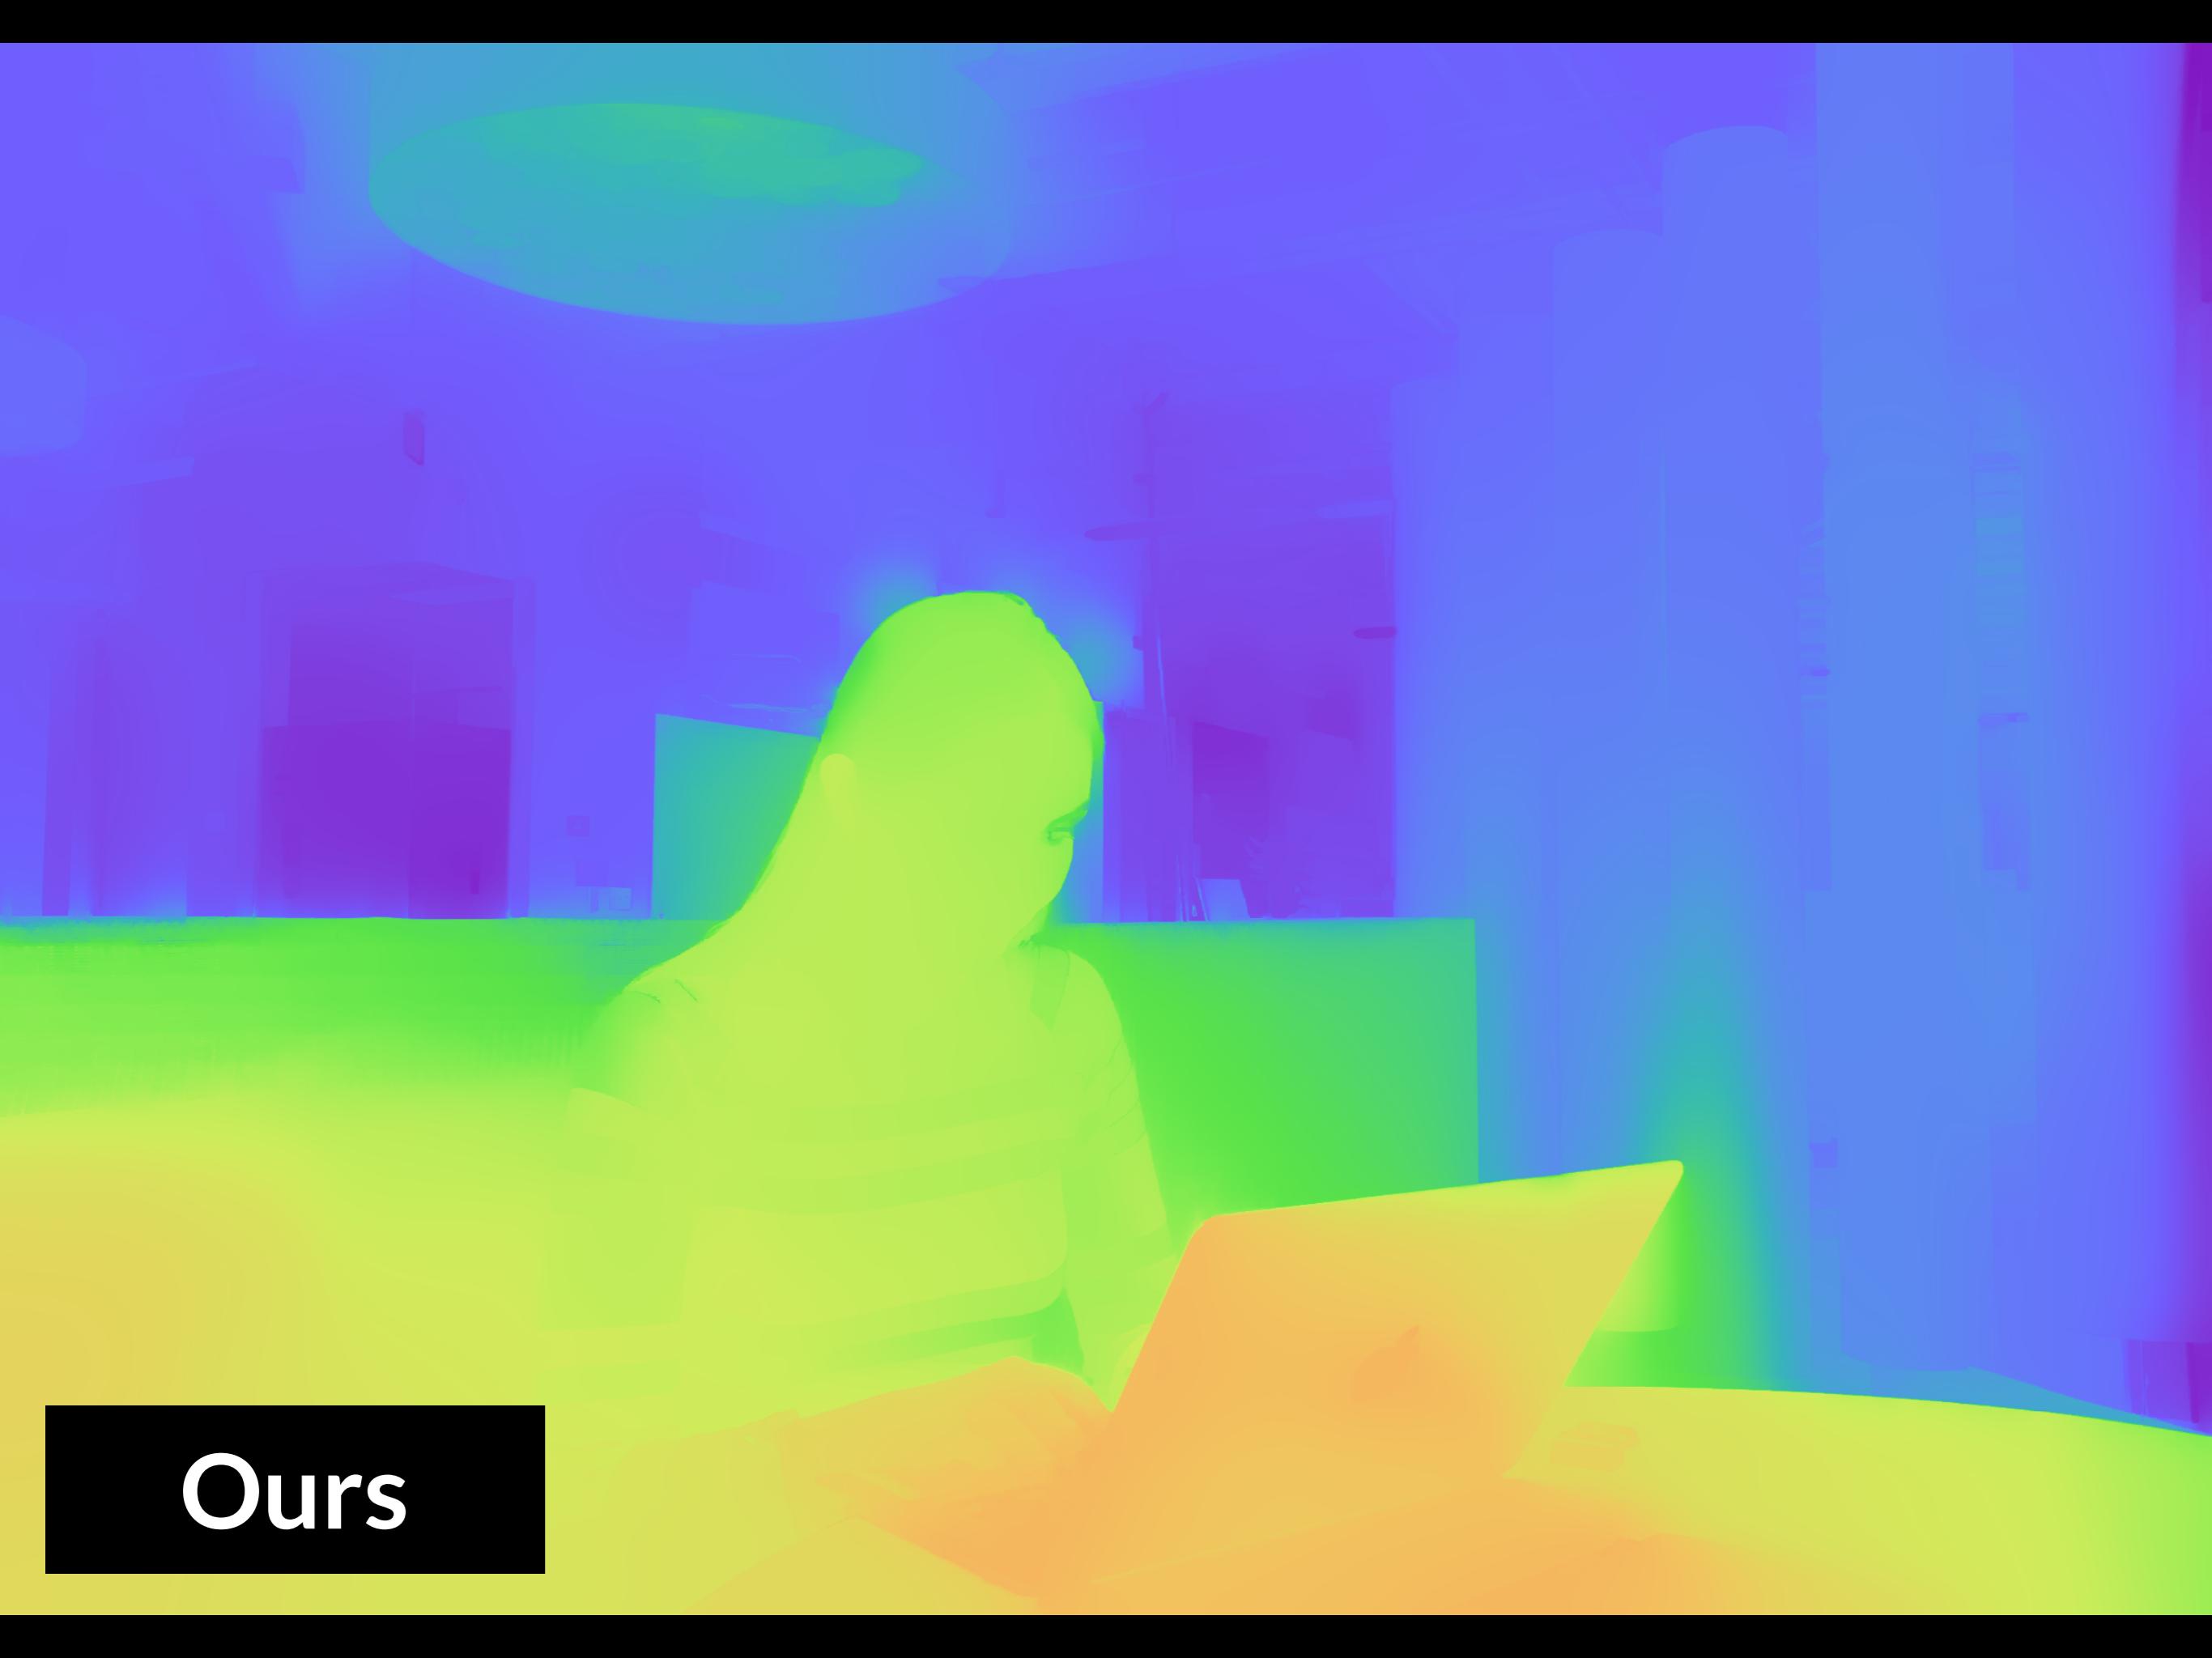
# Results



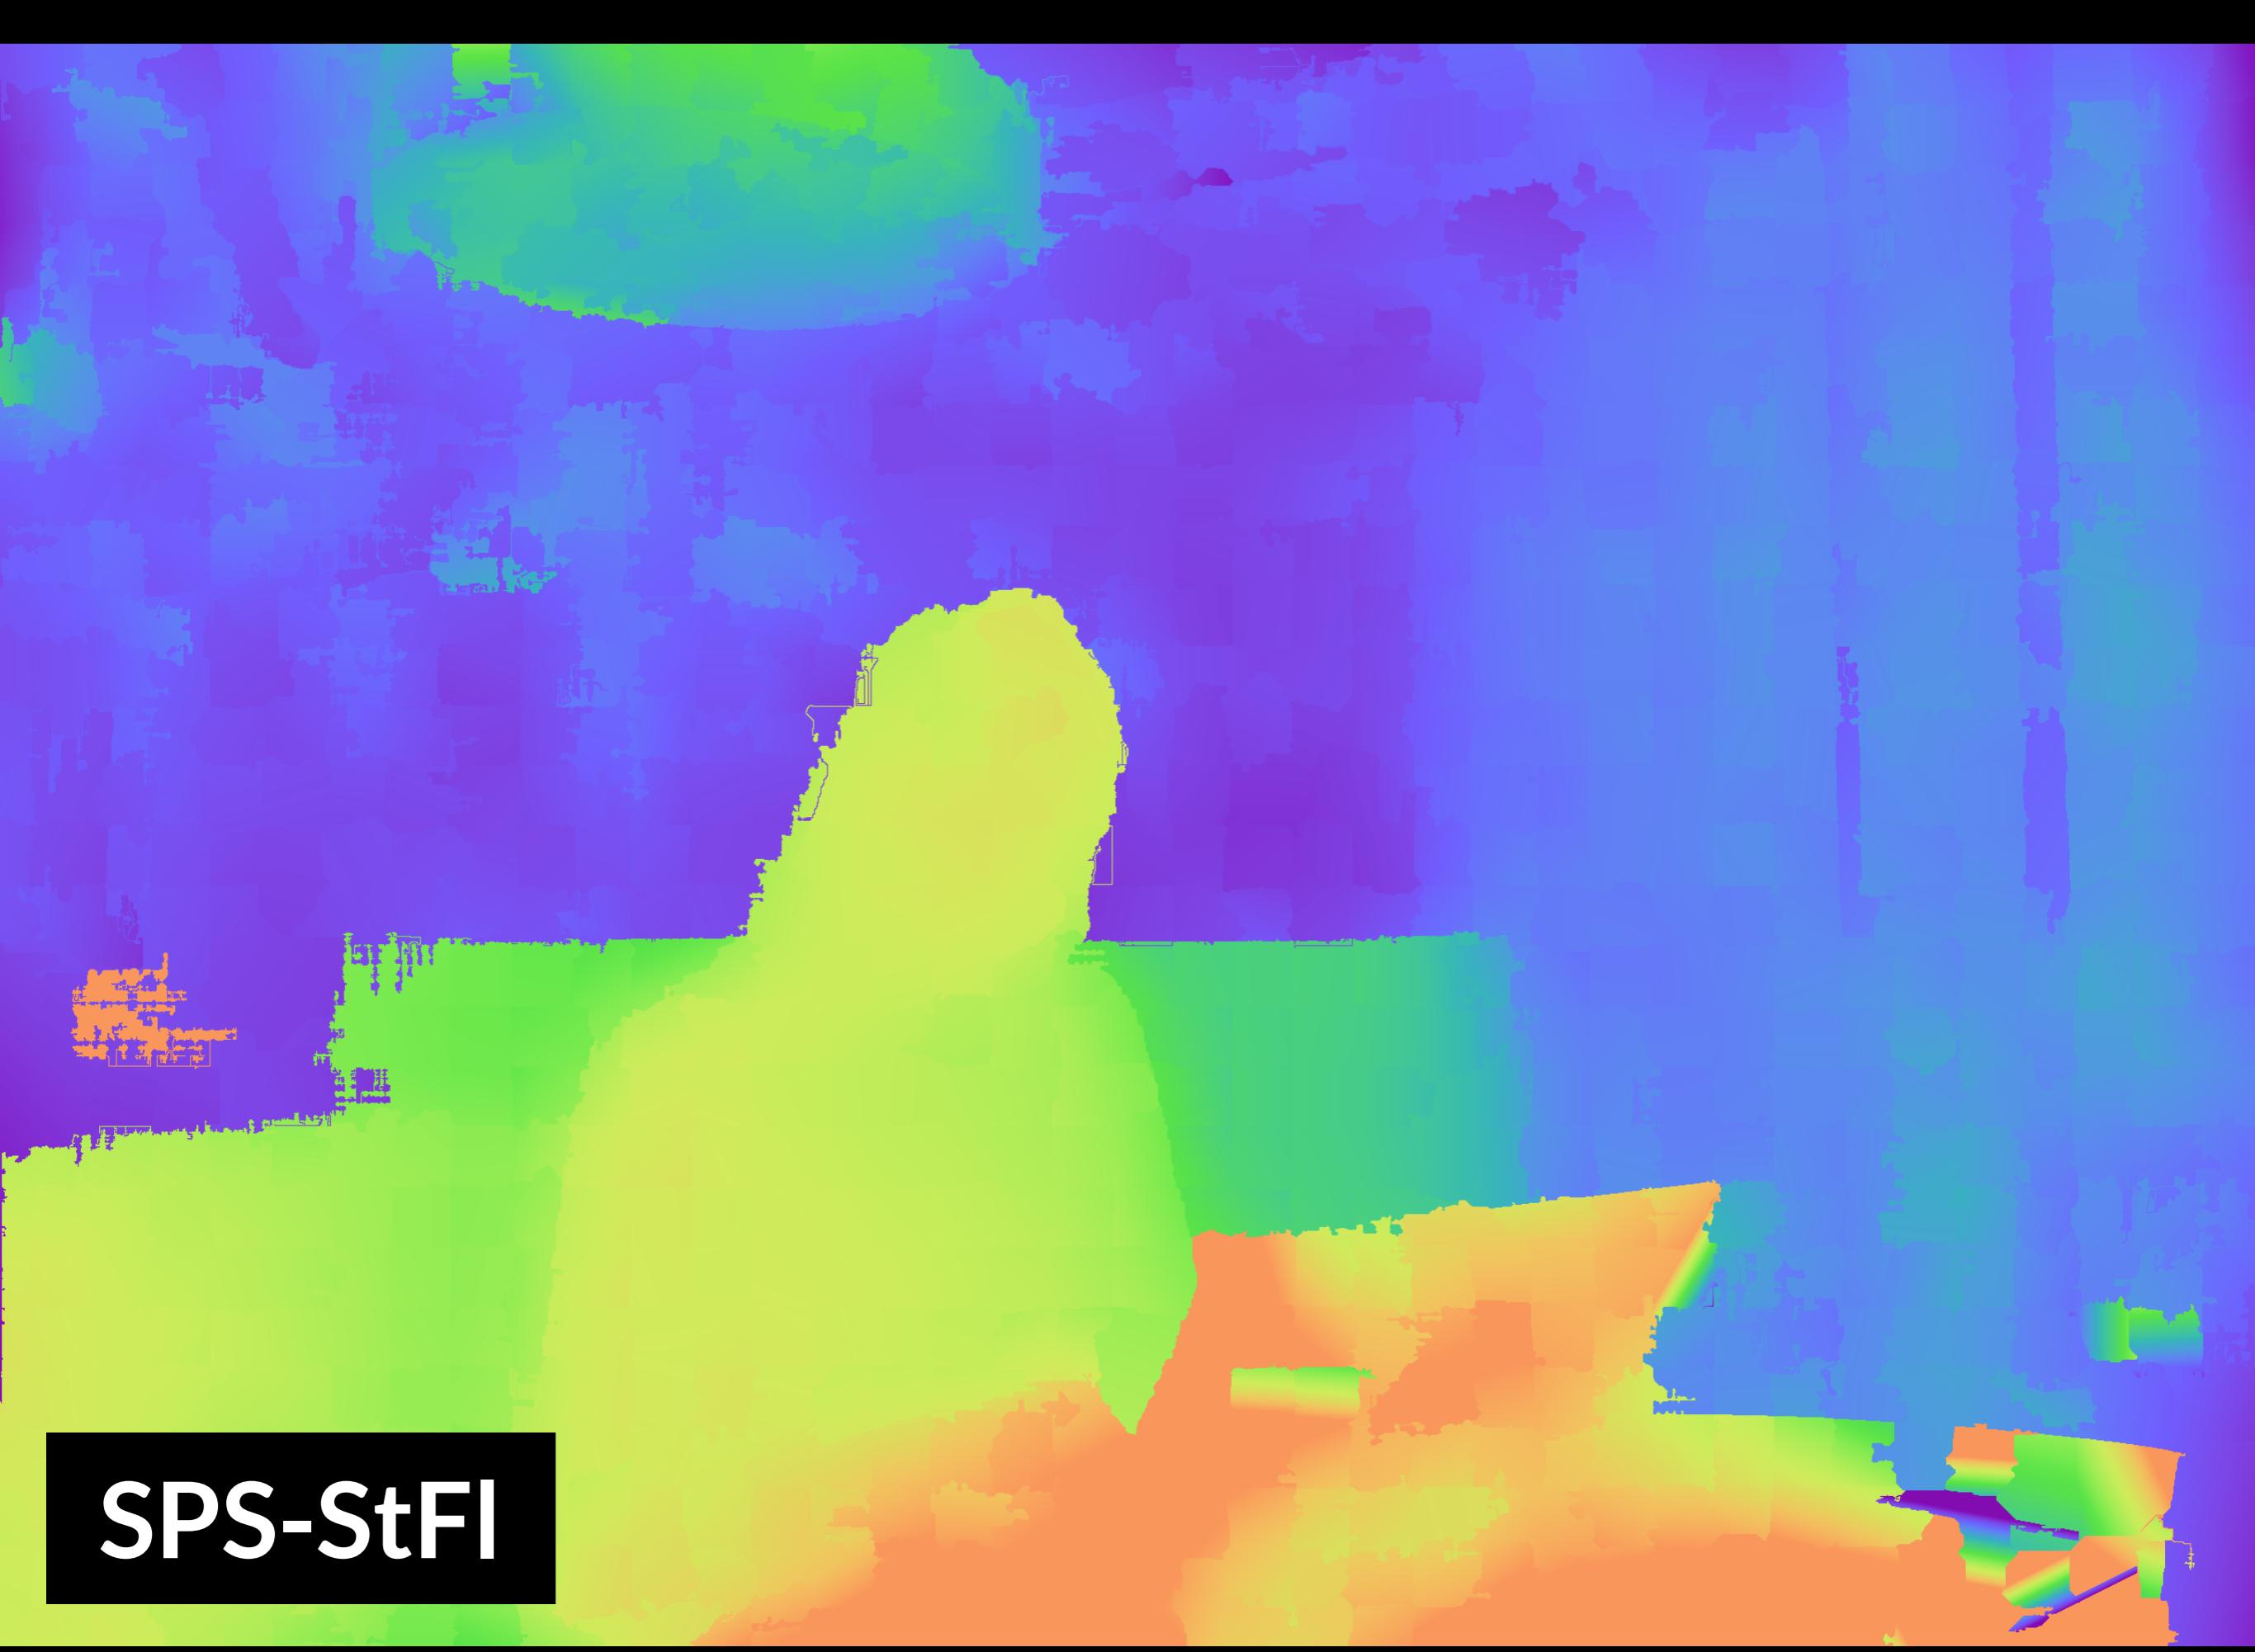
Input 1



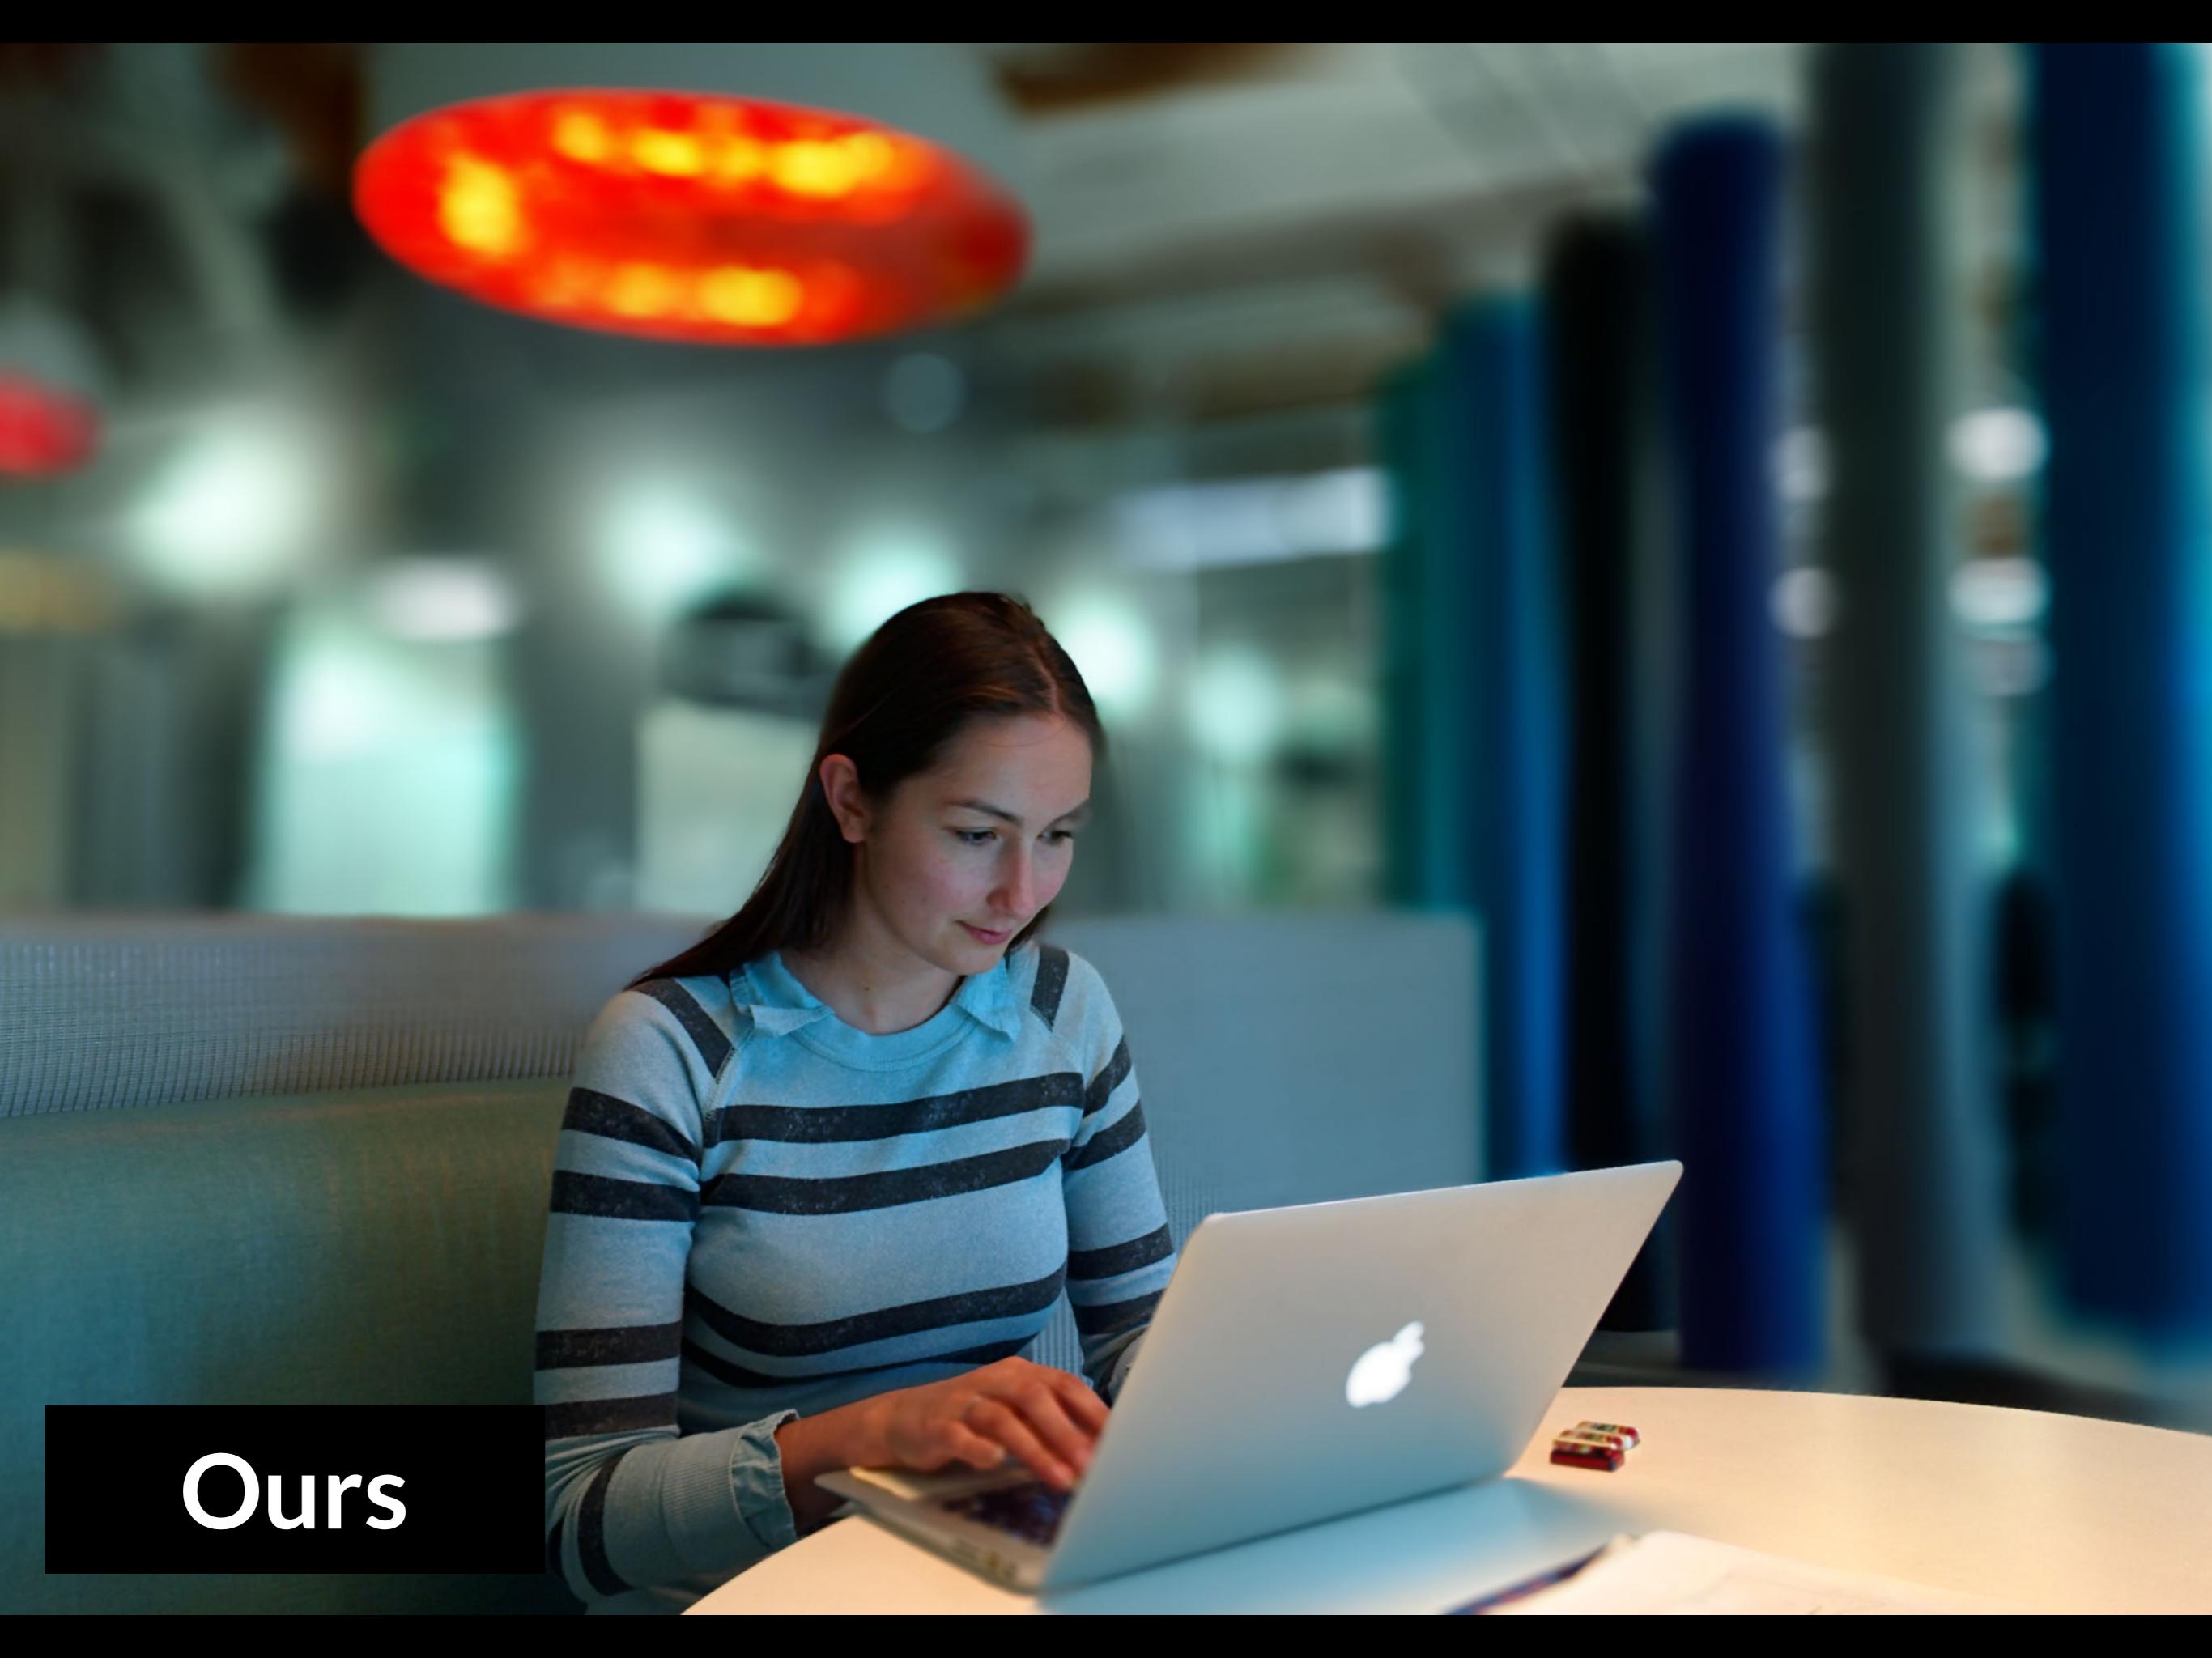
Input 2



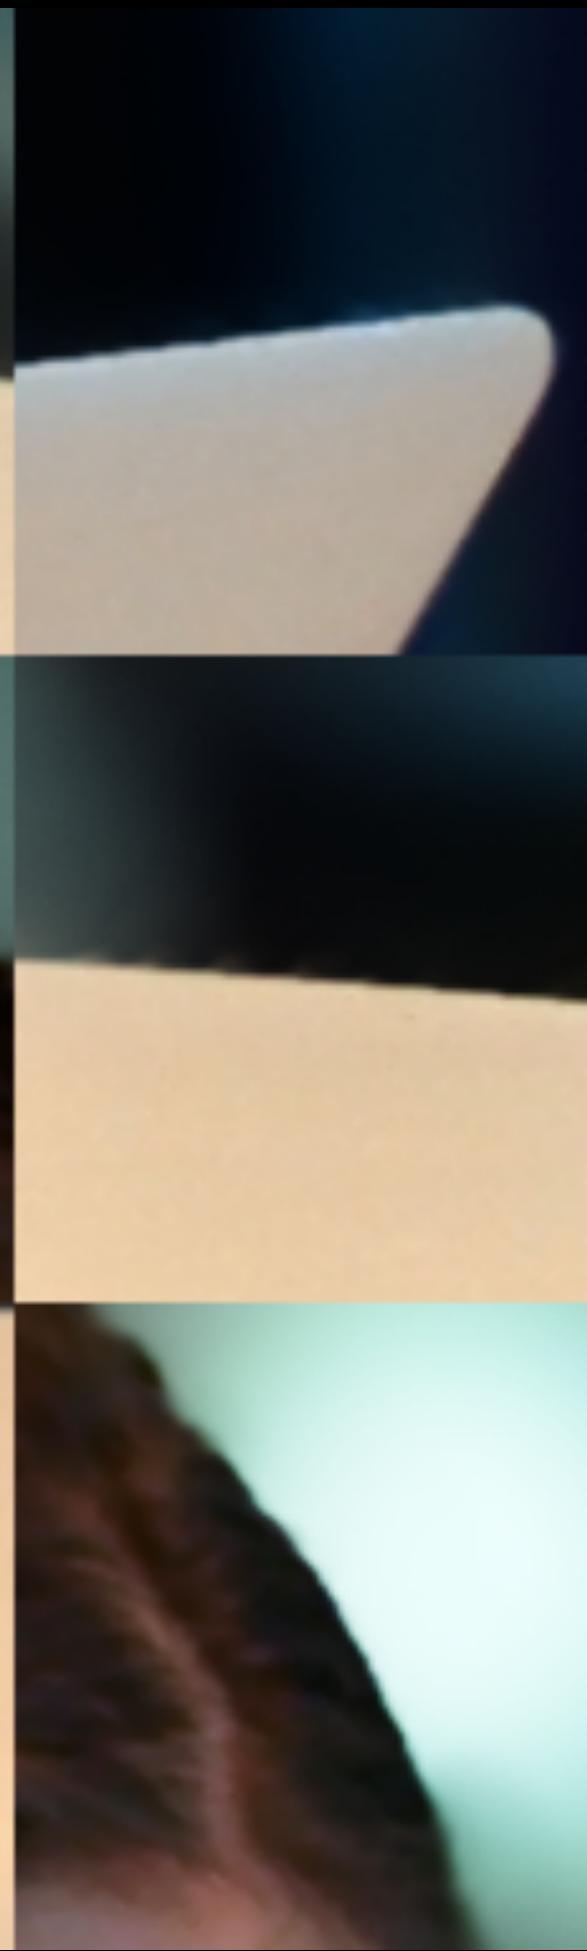
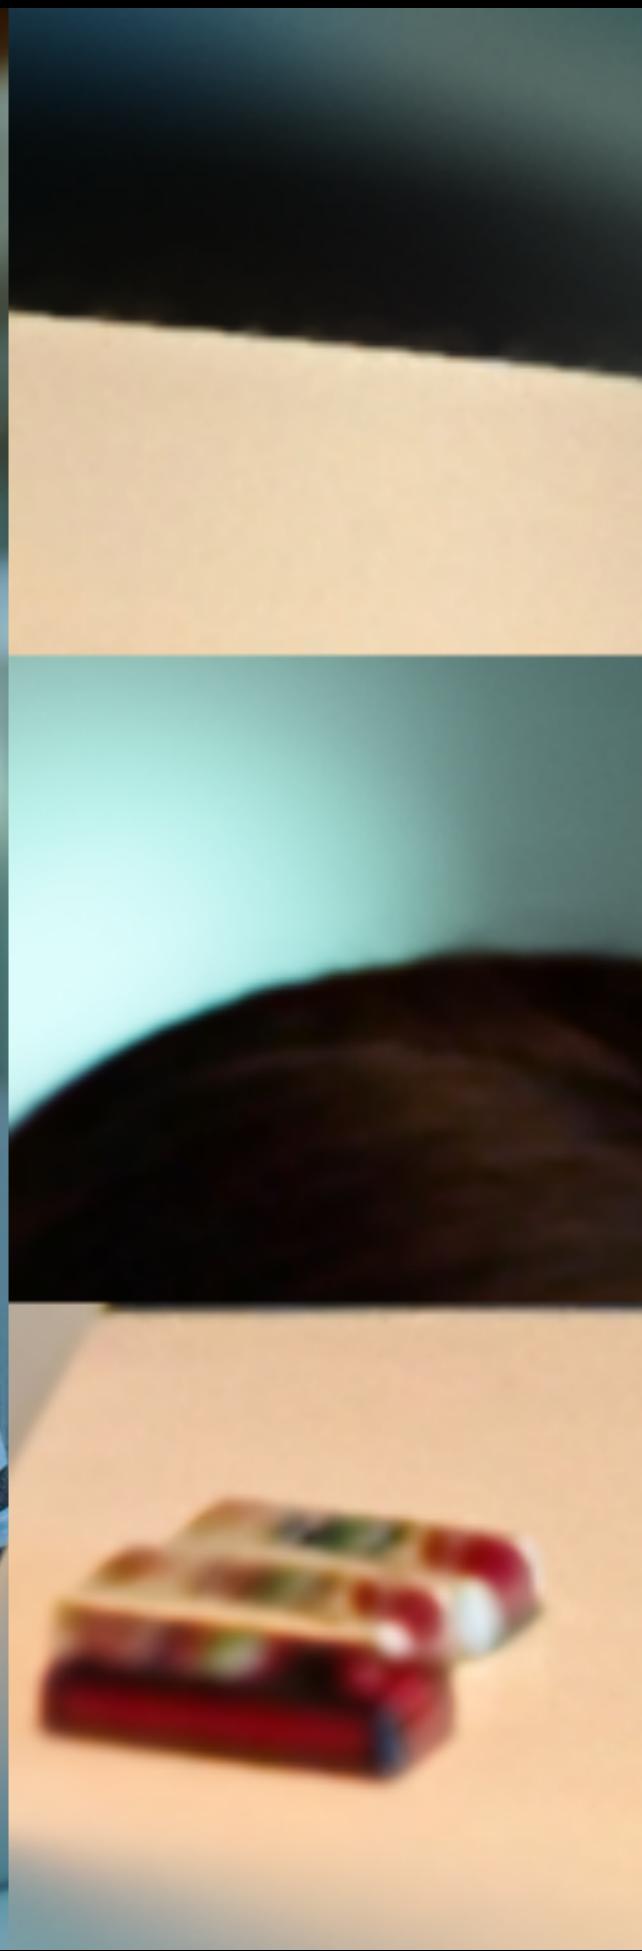
Ours



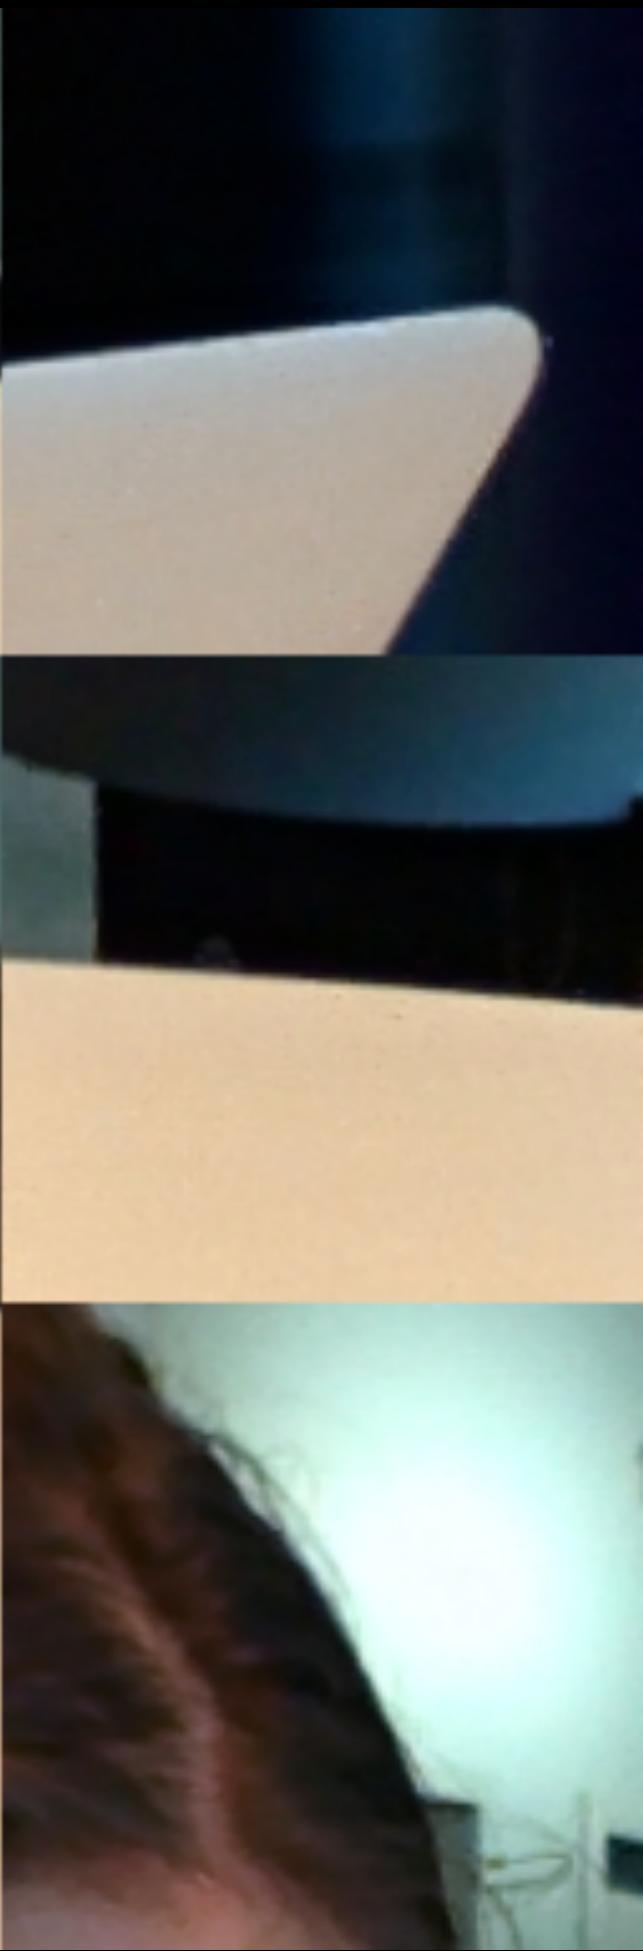
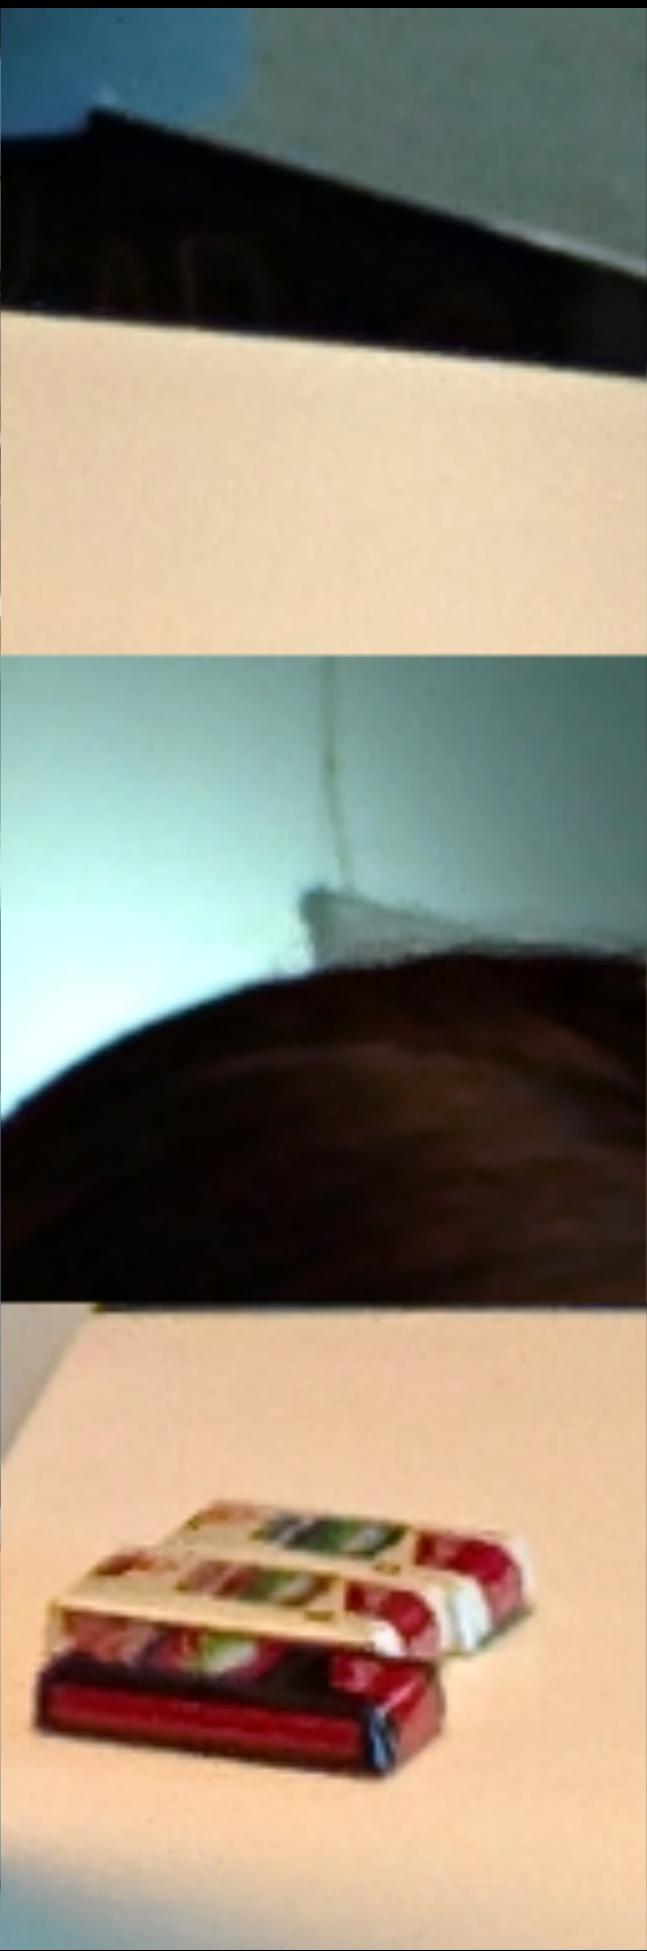
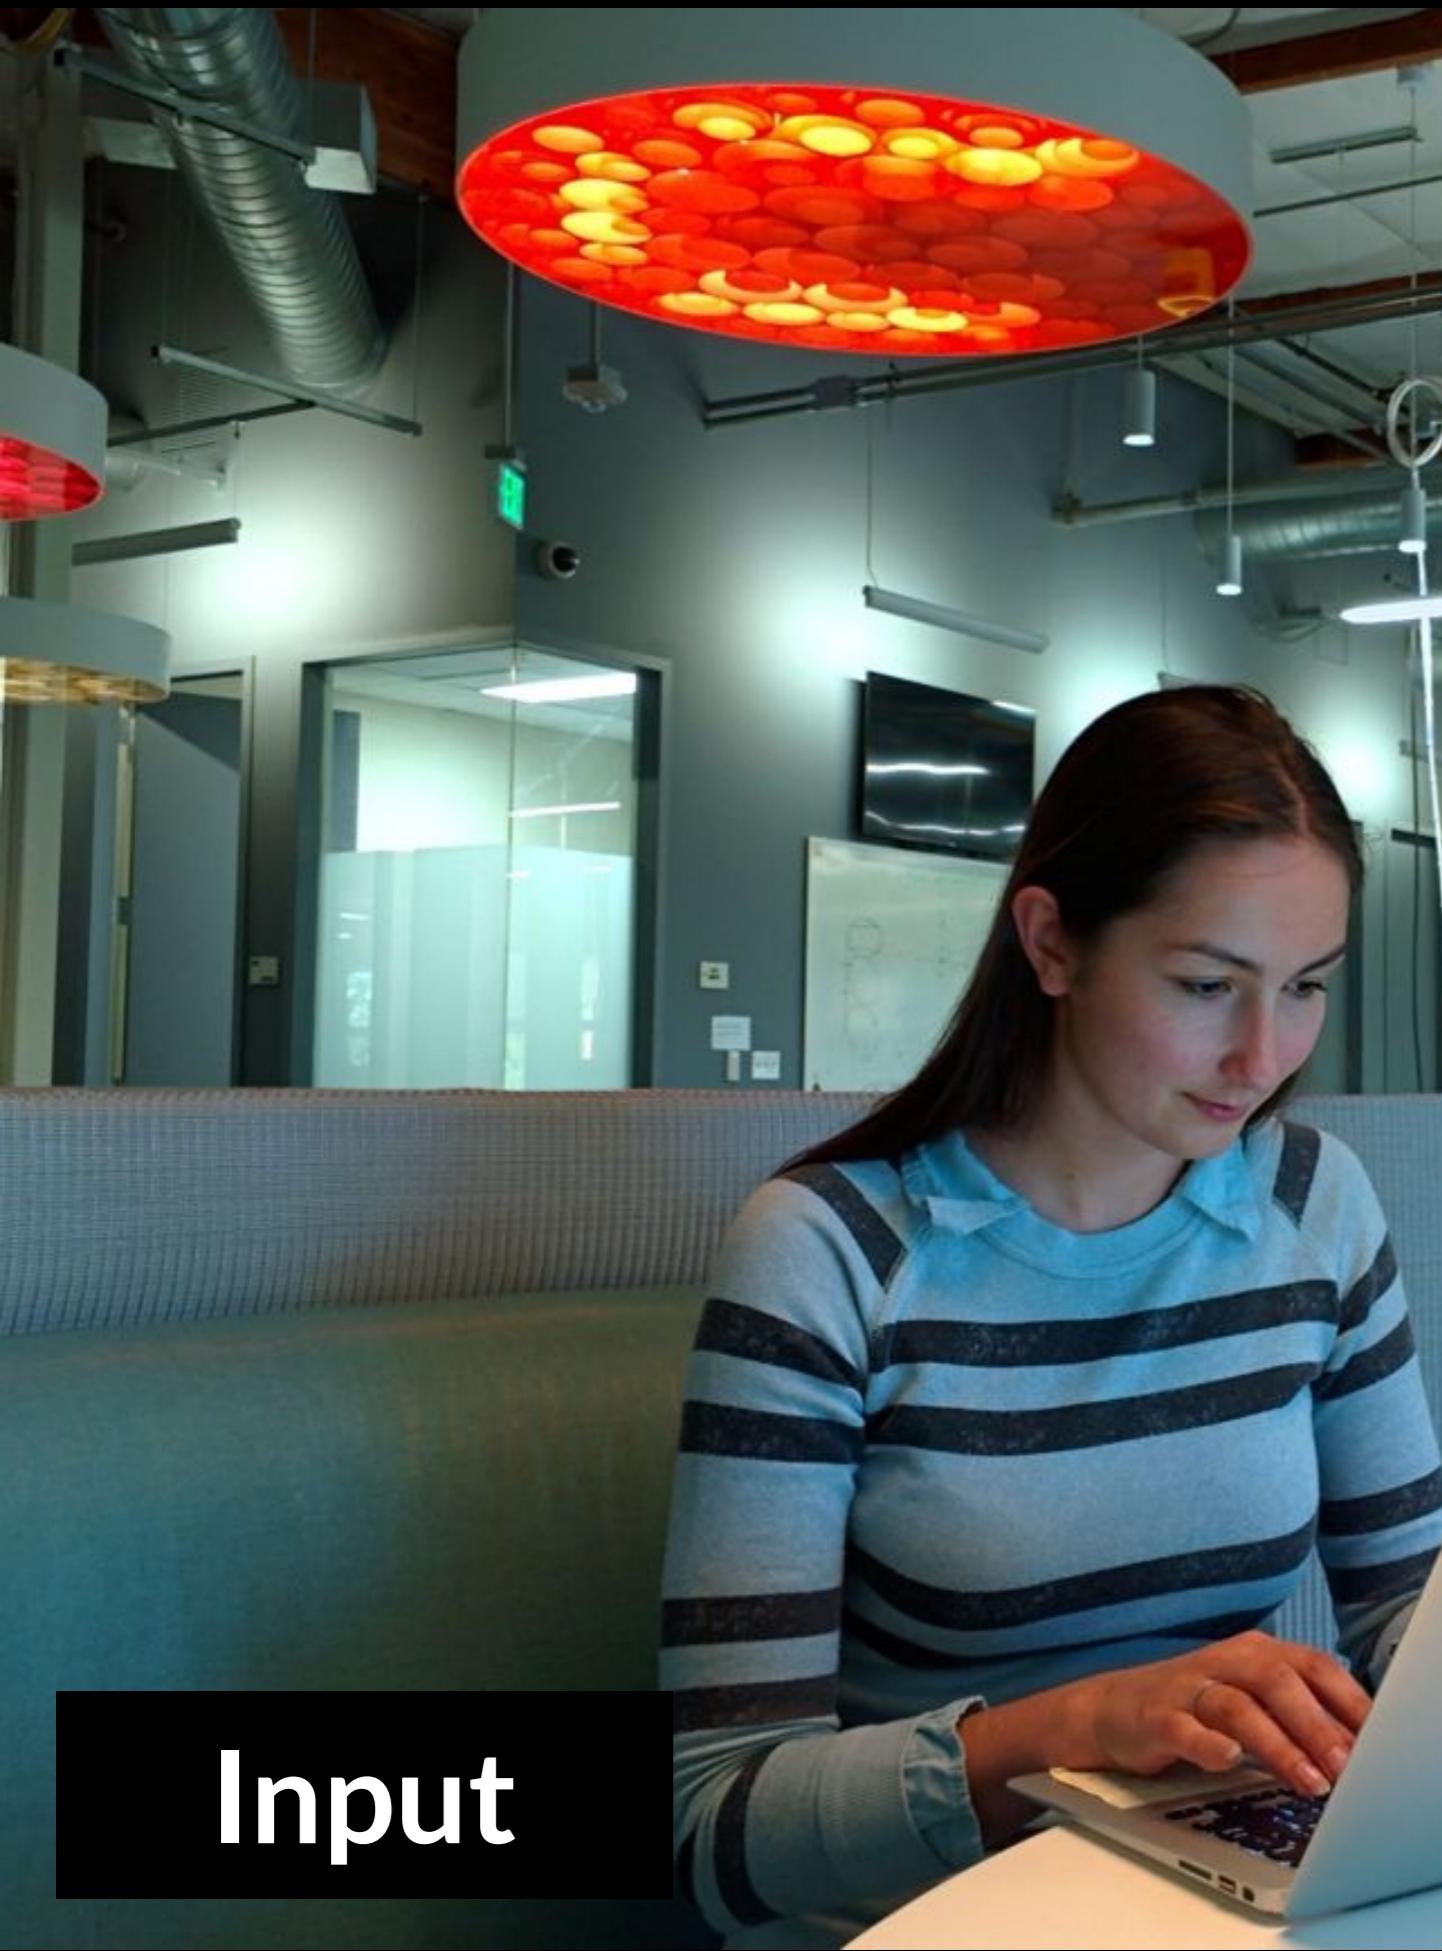
Ours



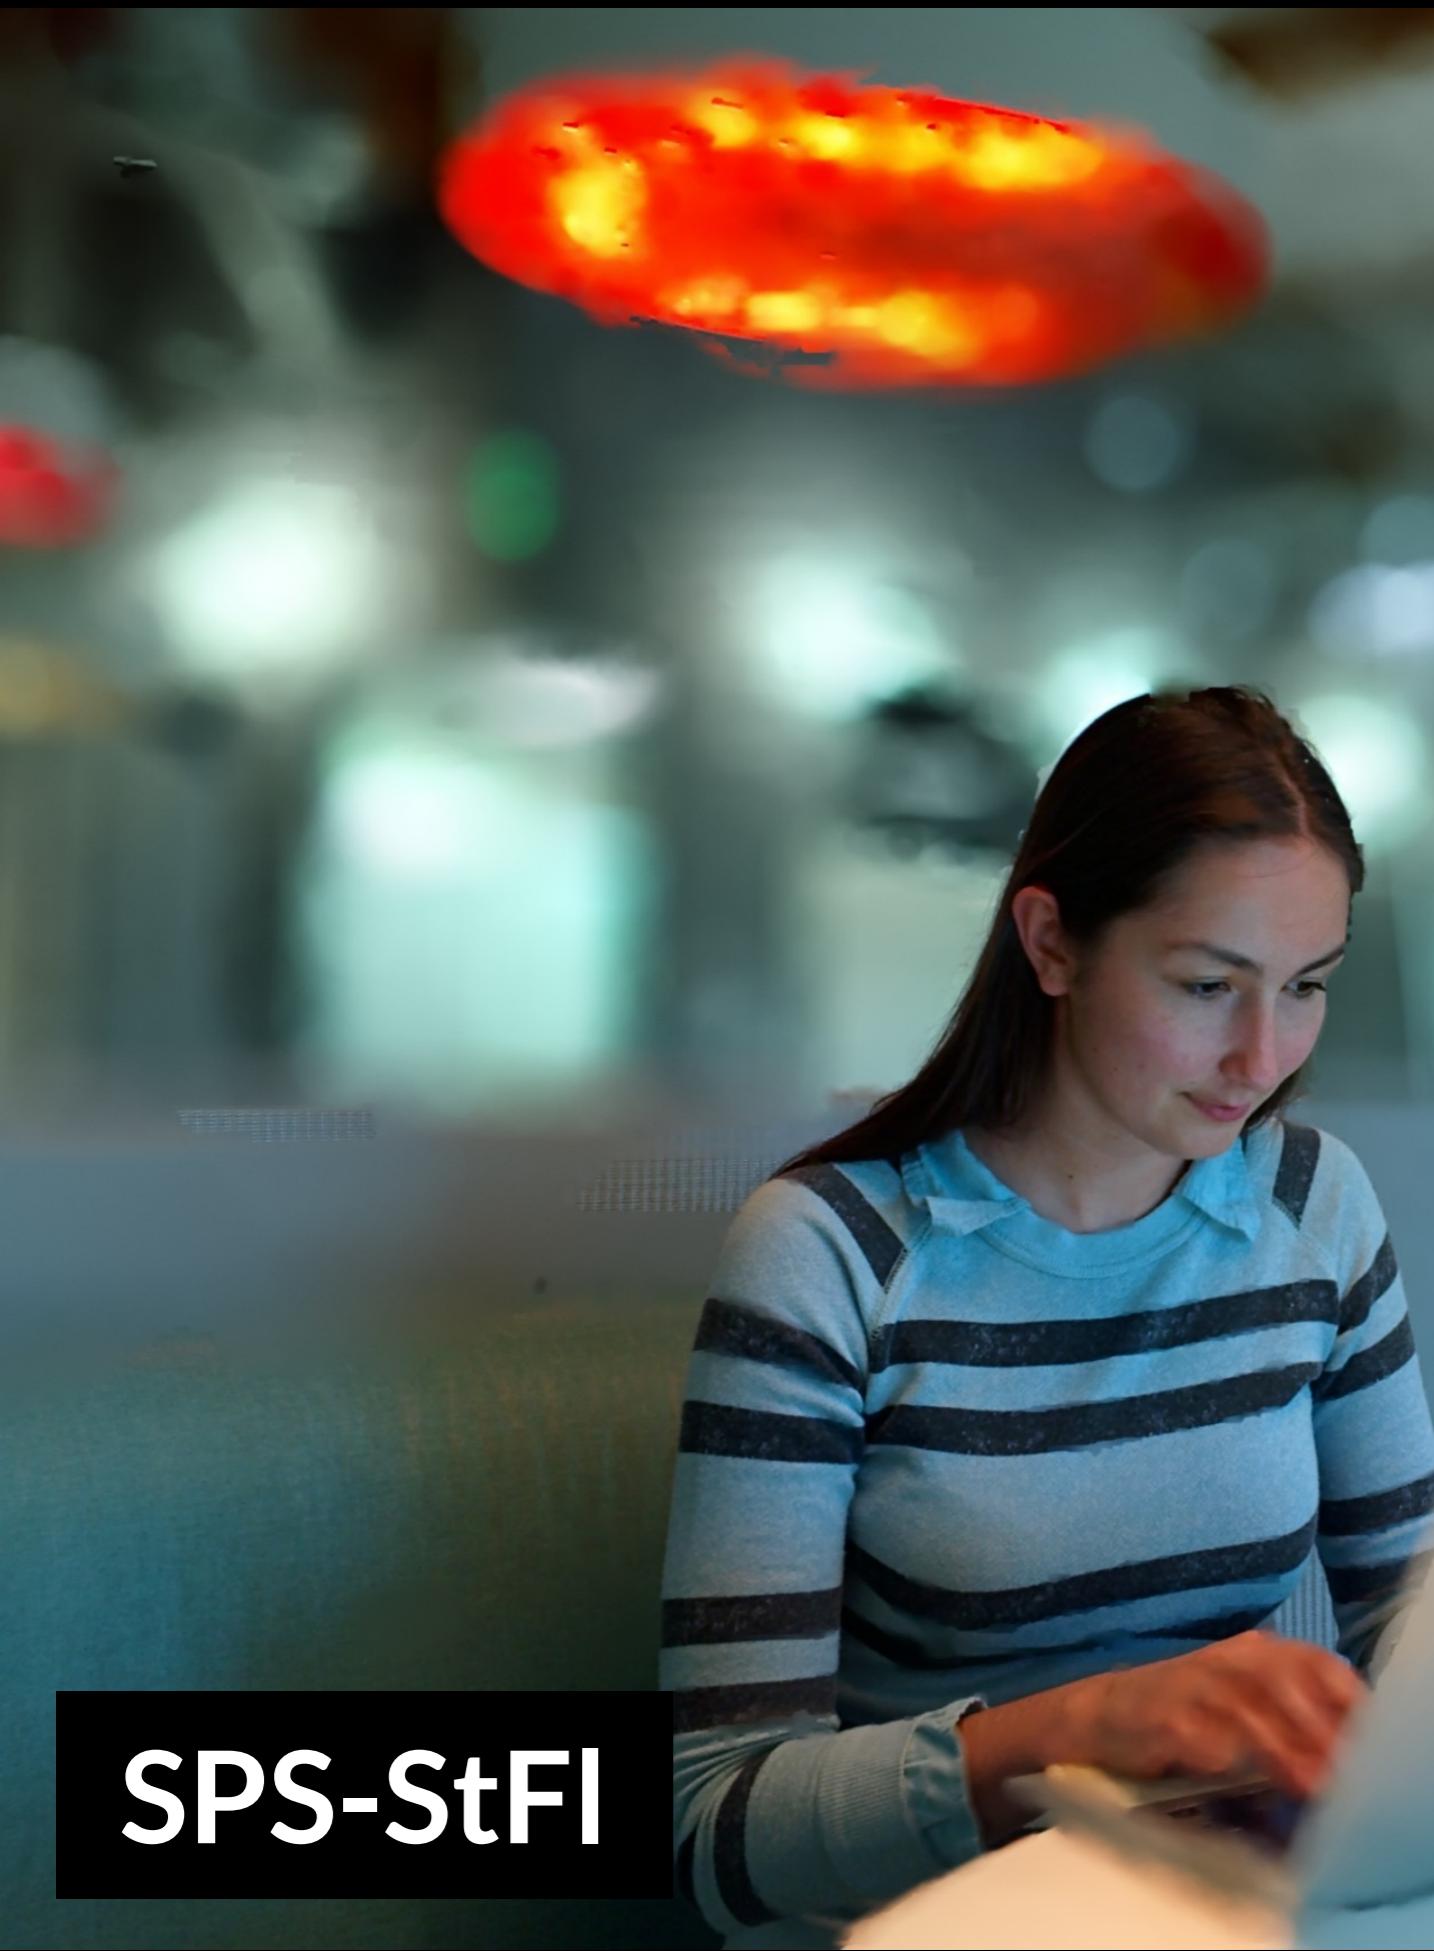
Ours



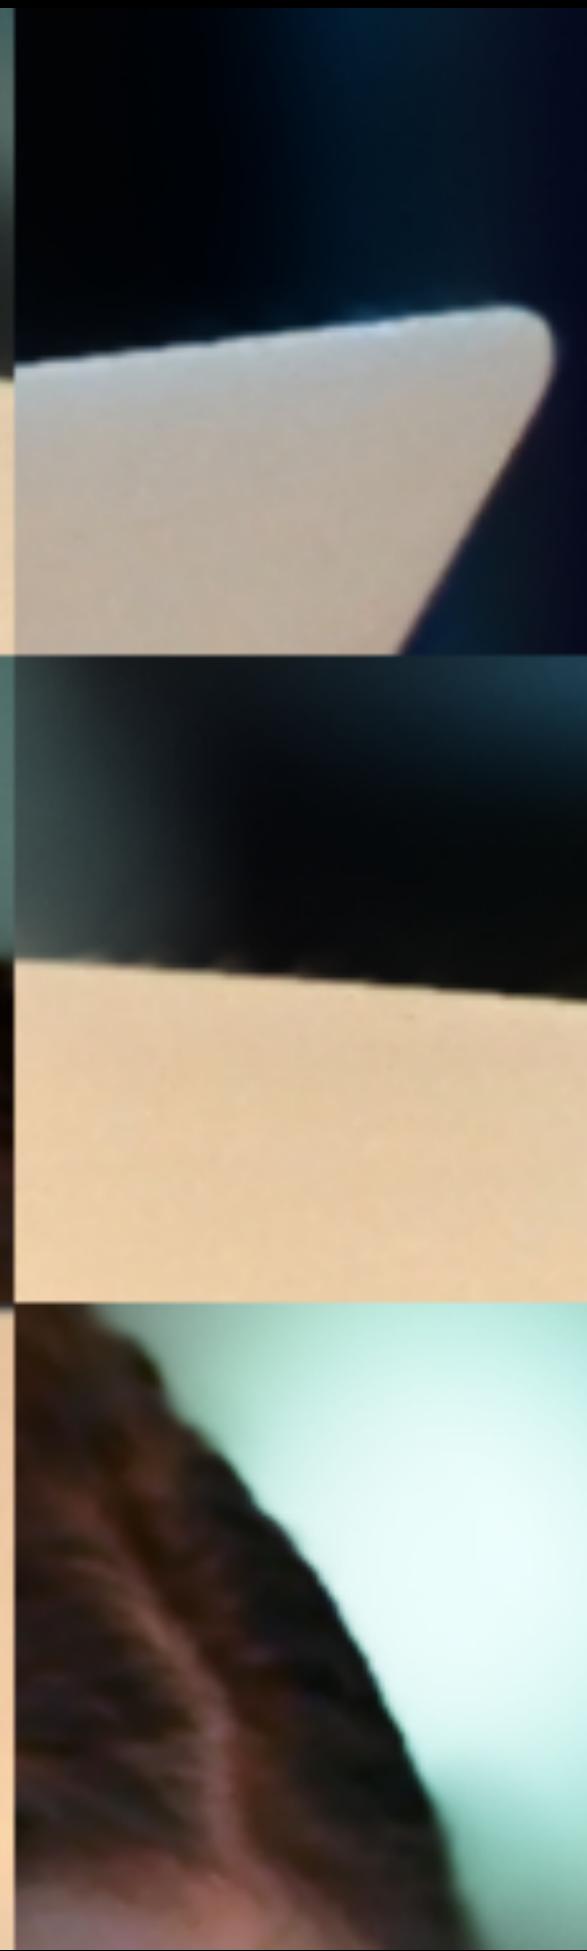
**Input**



**SPS-StFI**



Ours



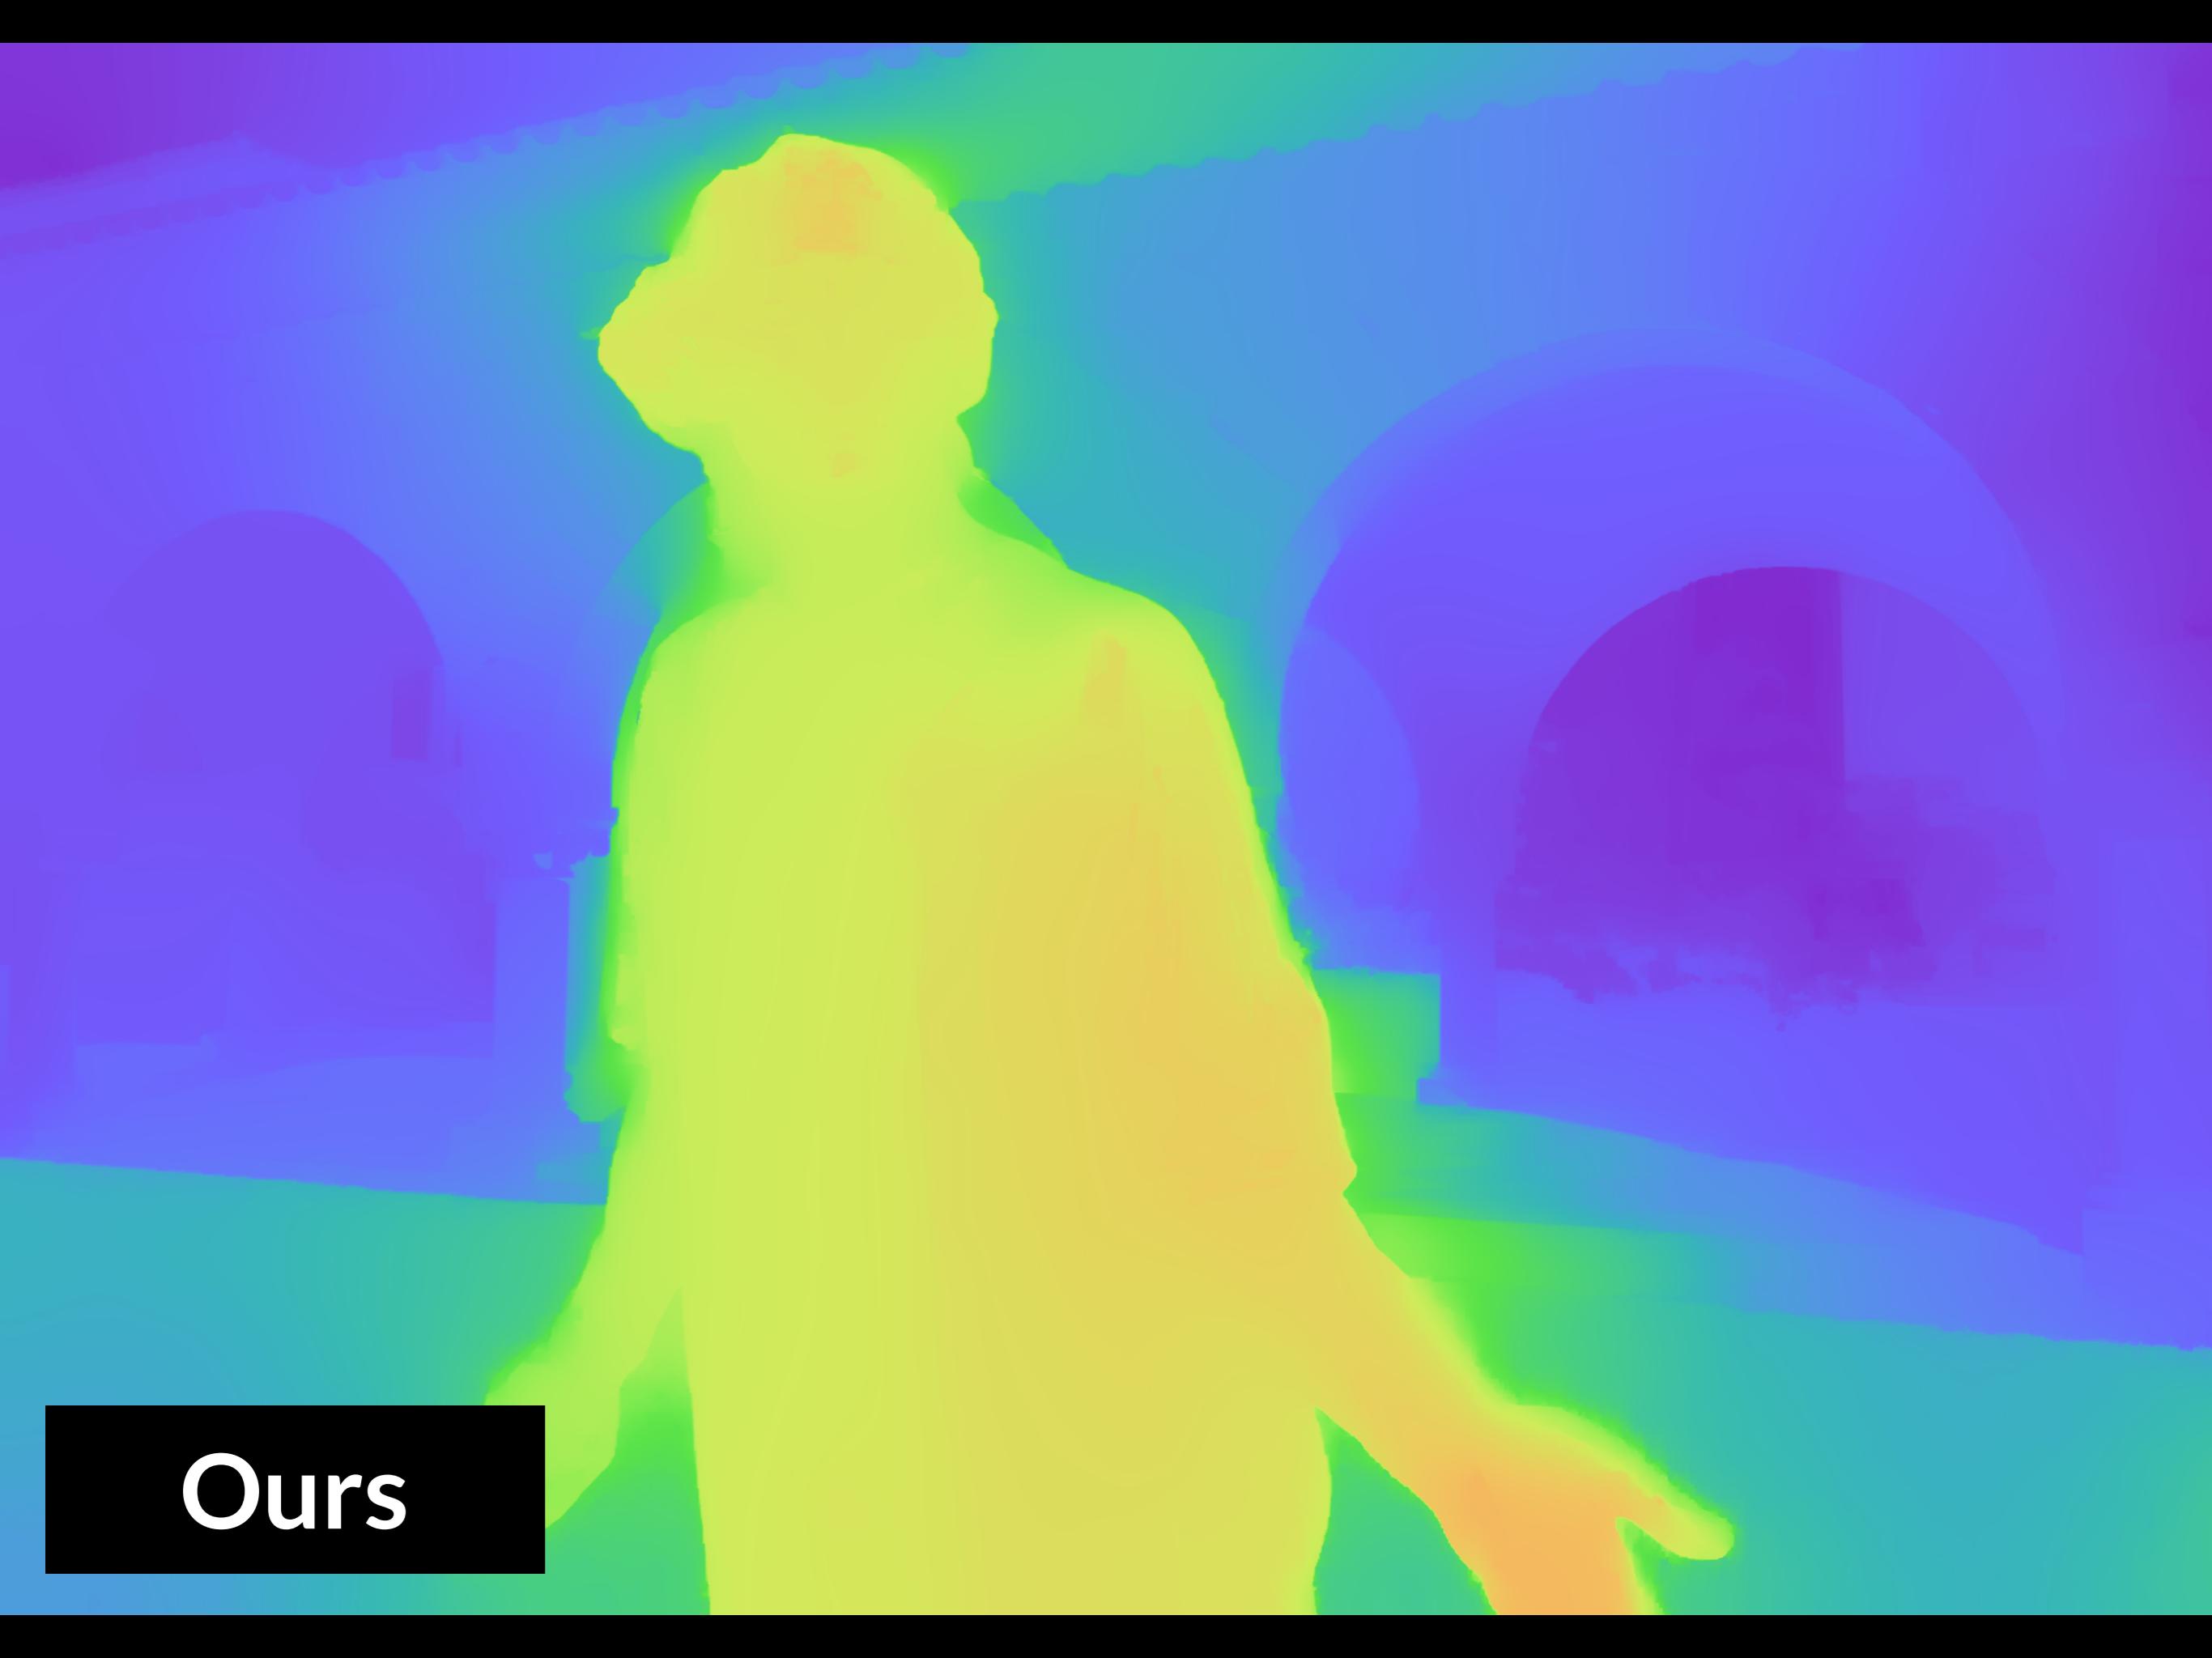


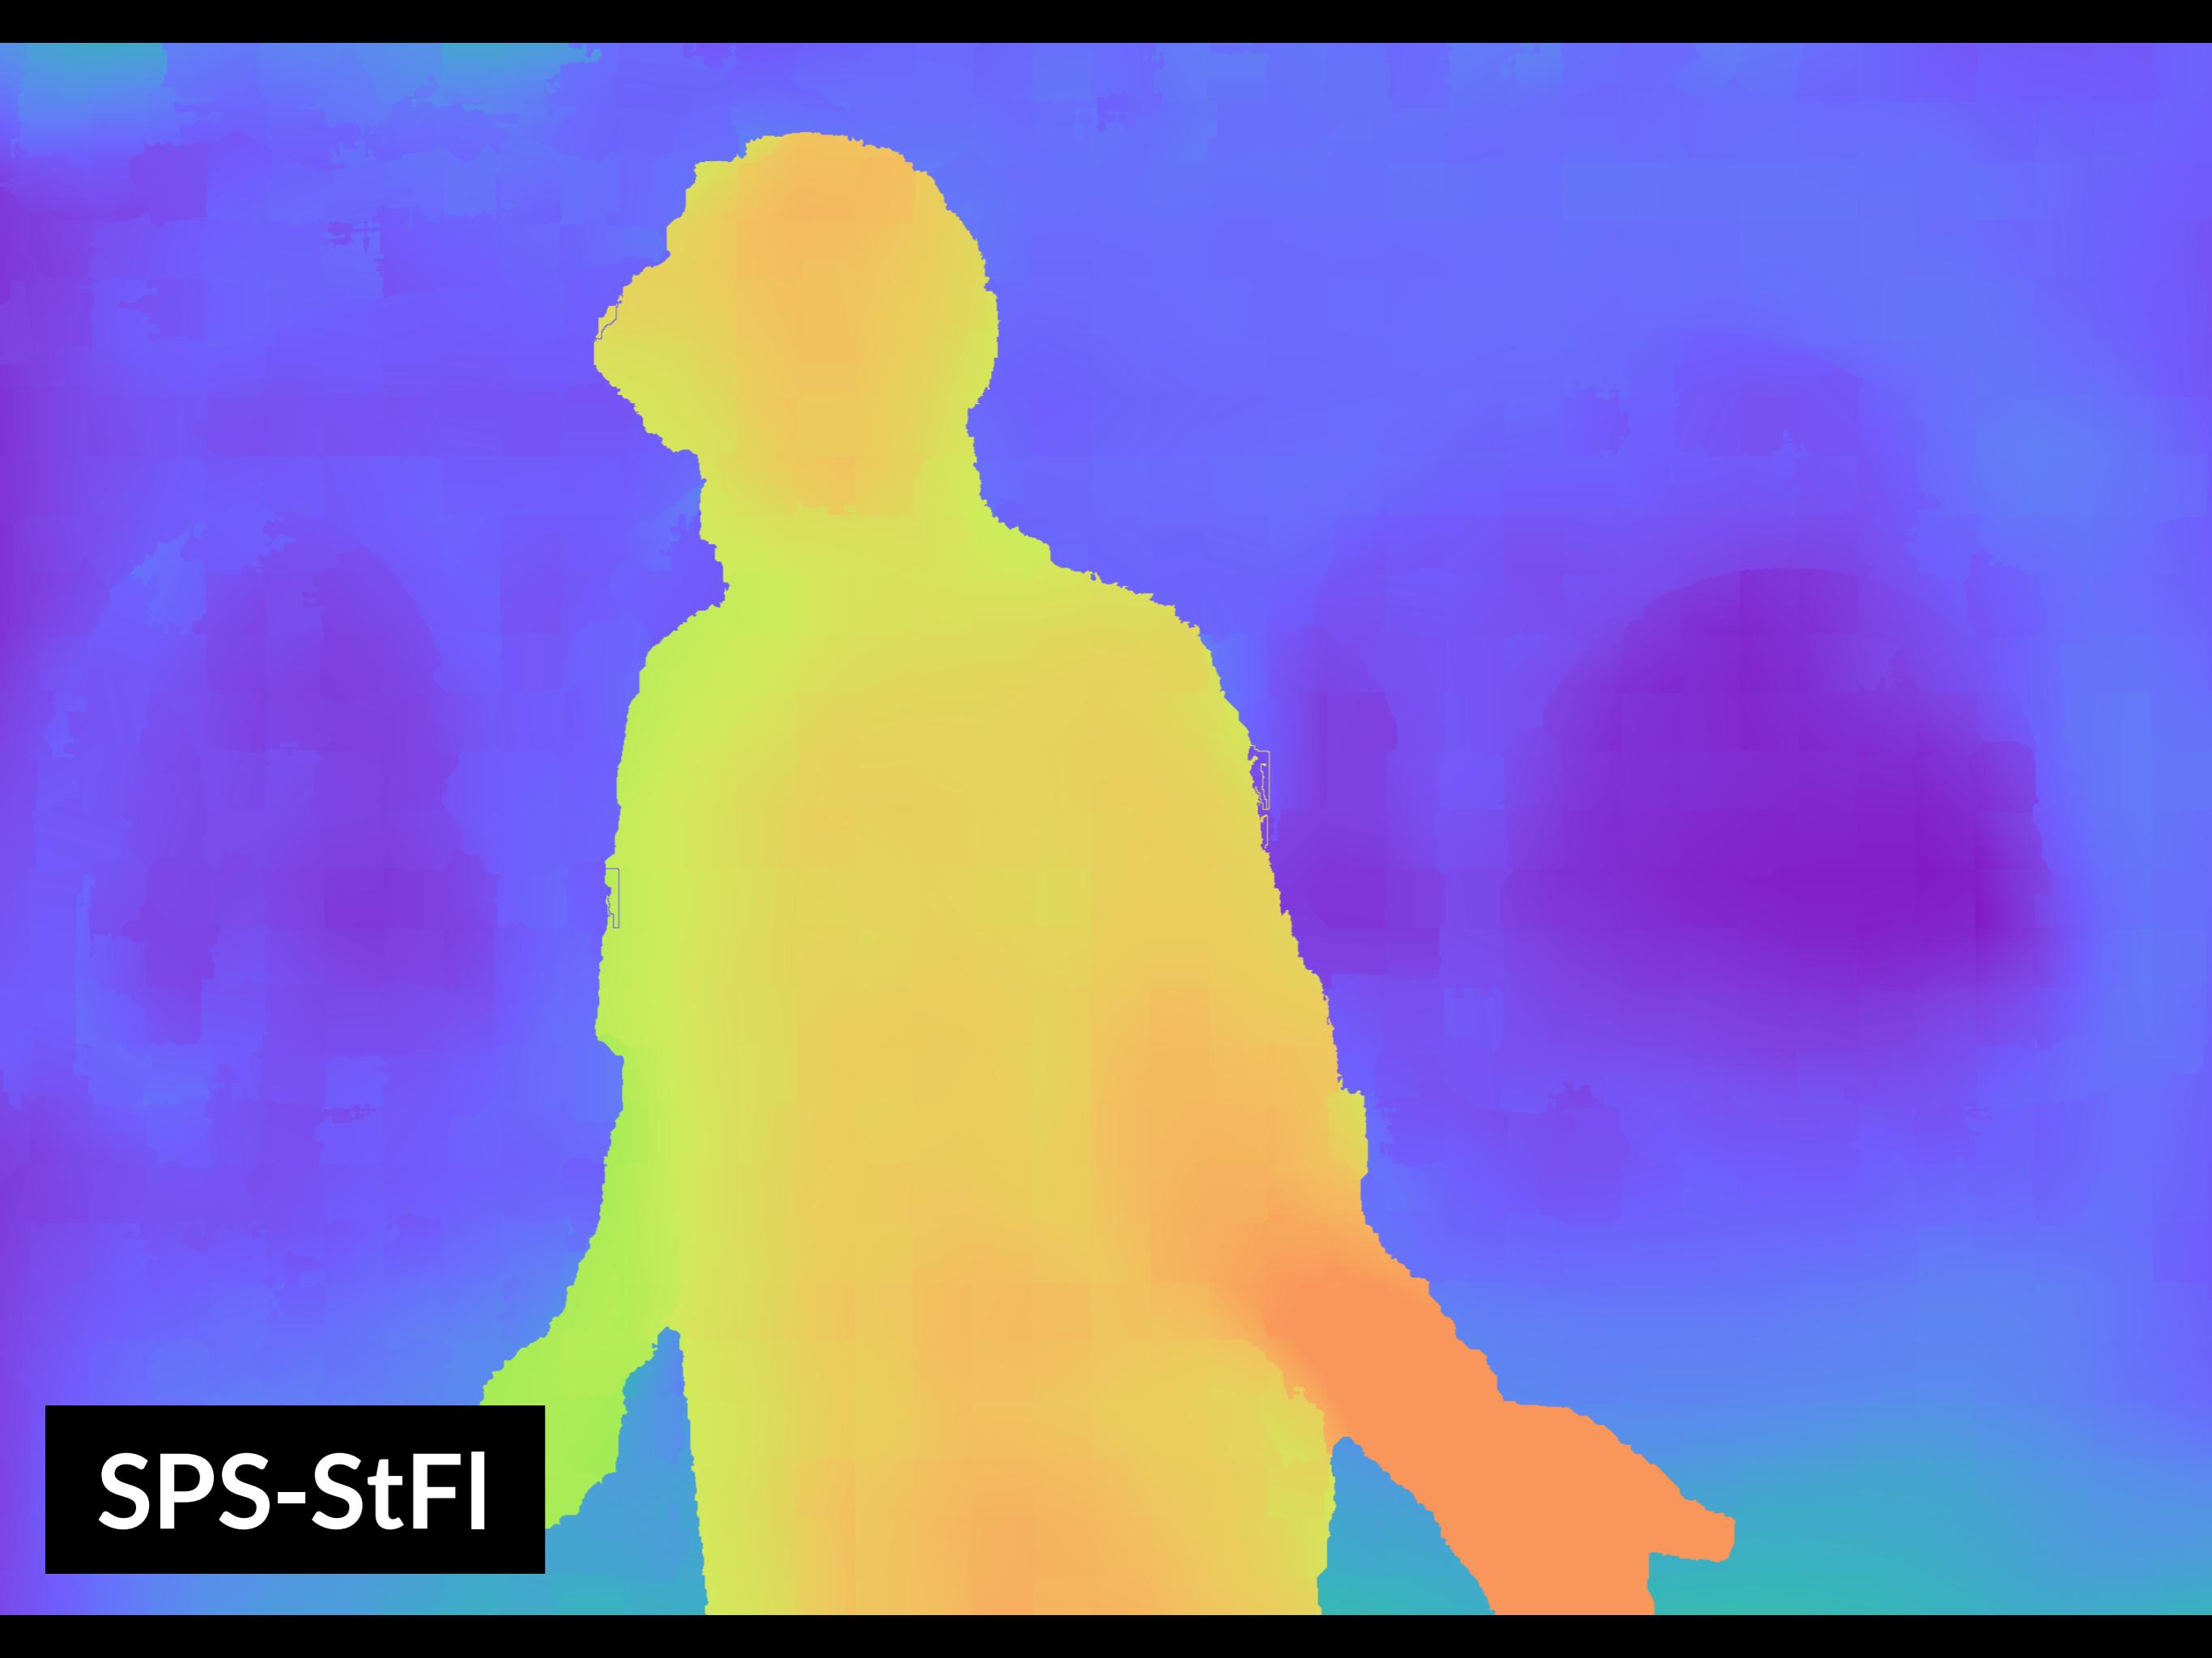
Input 1

**Input 2**



Ours



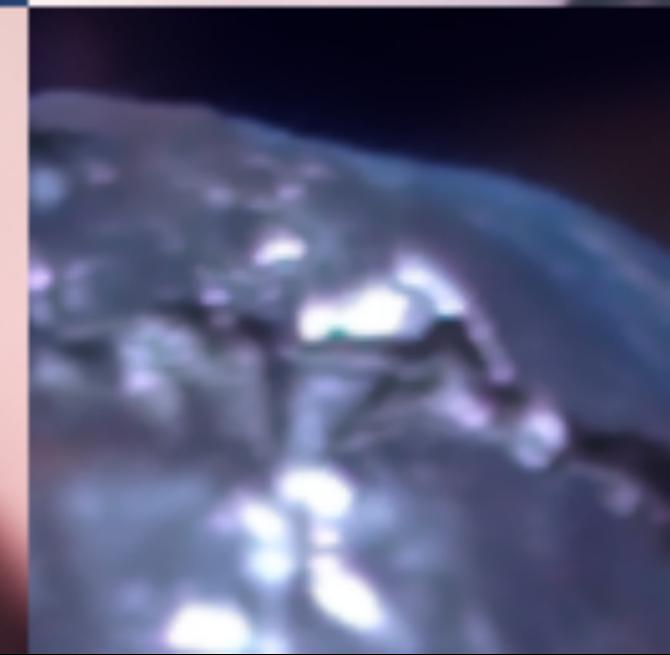


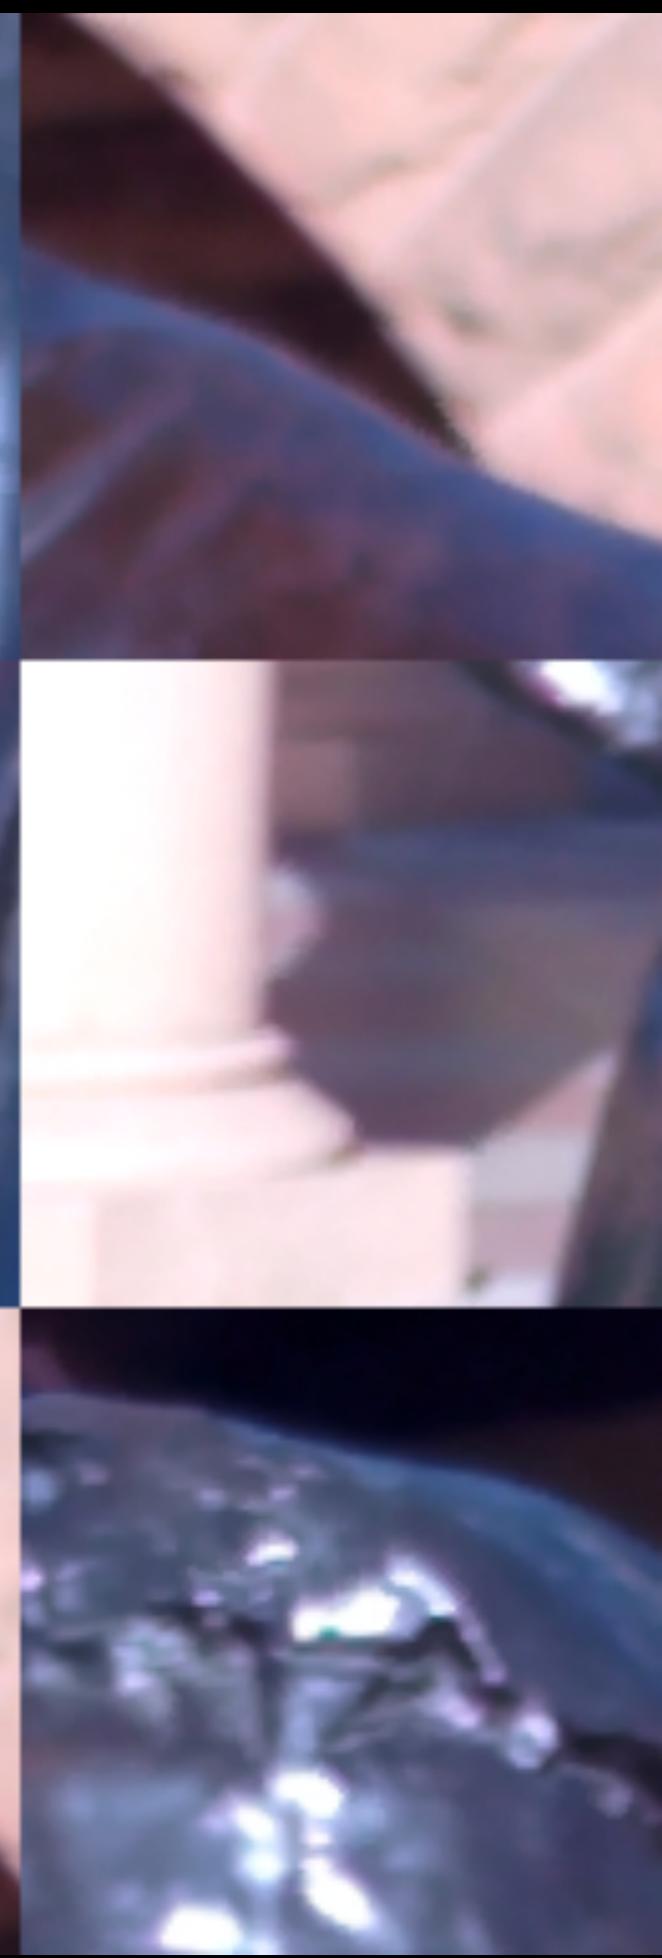
**SPS-StFI**



Ours

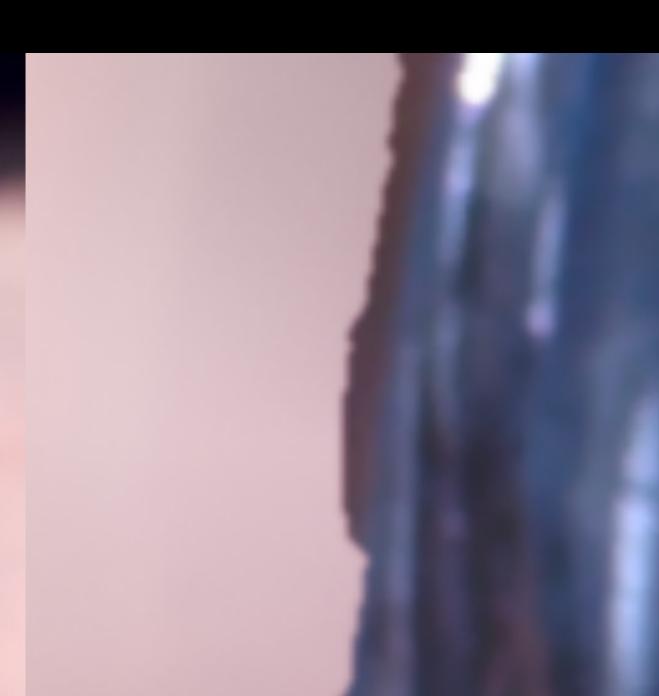
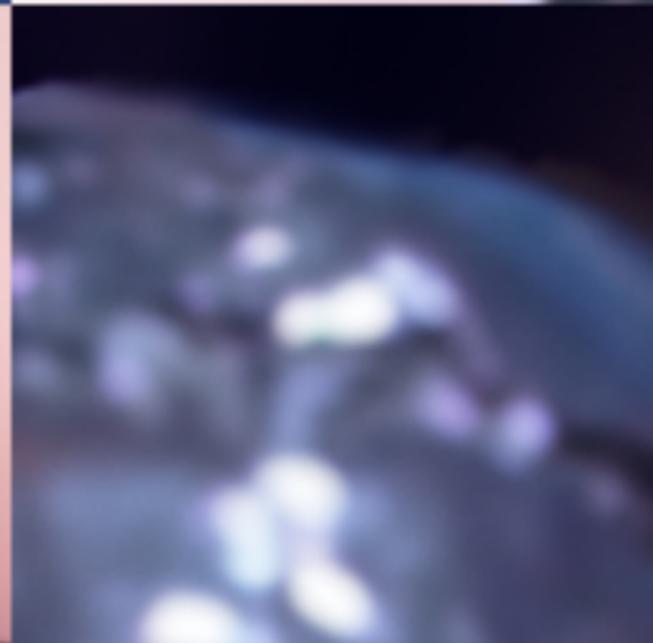
Ours



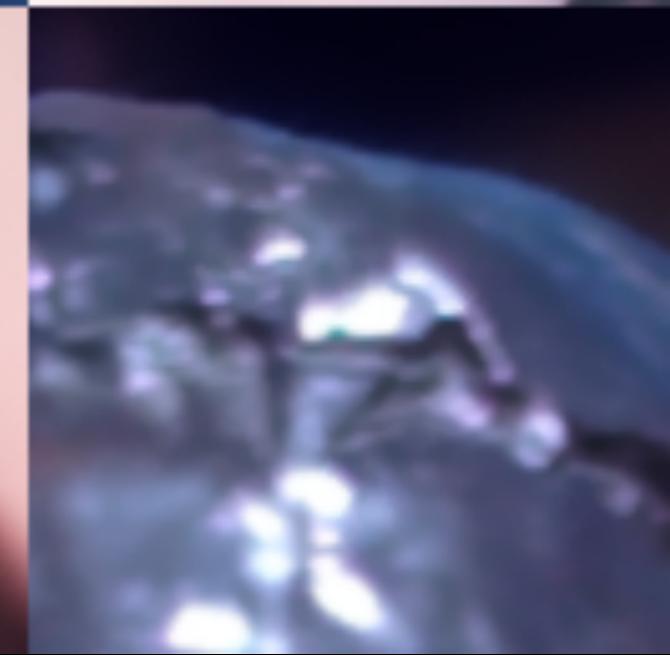




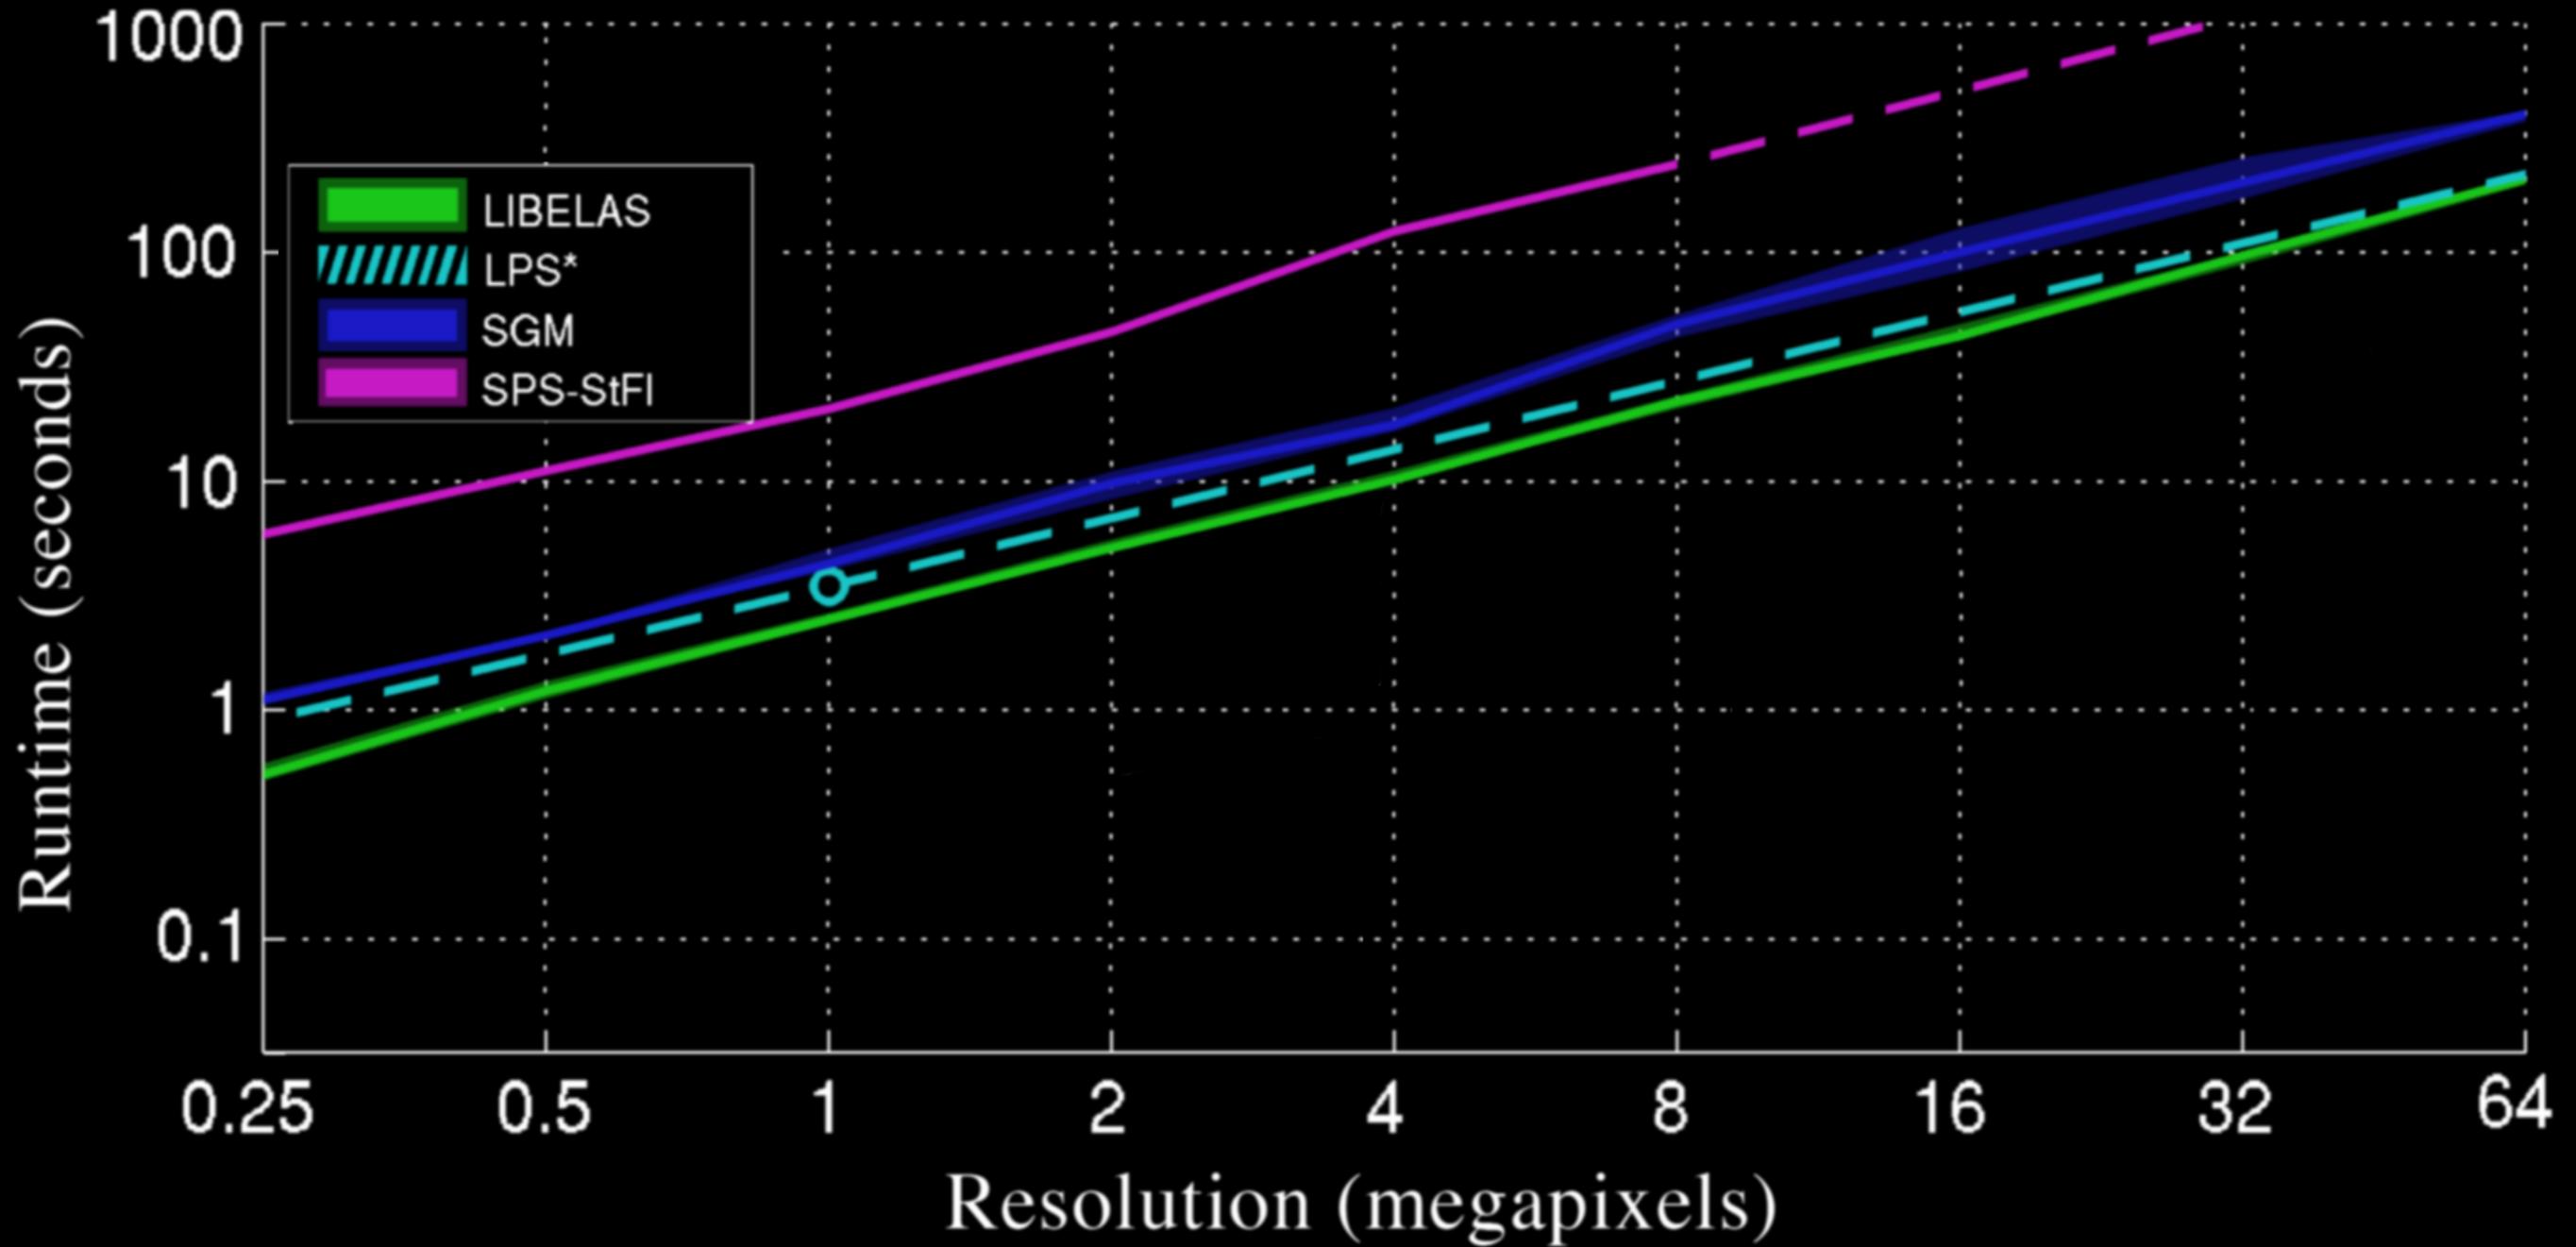
**SPS-StFI**



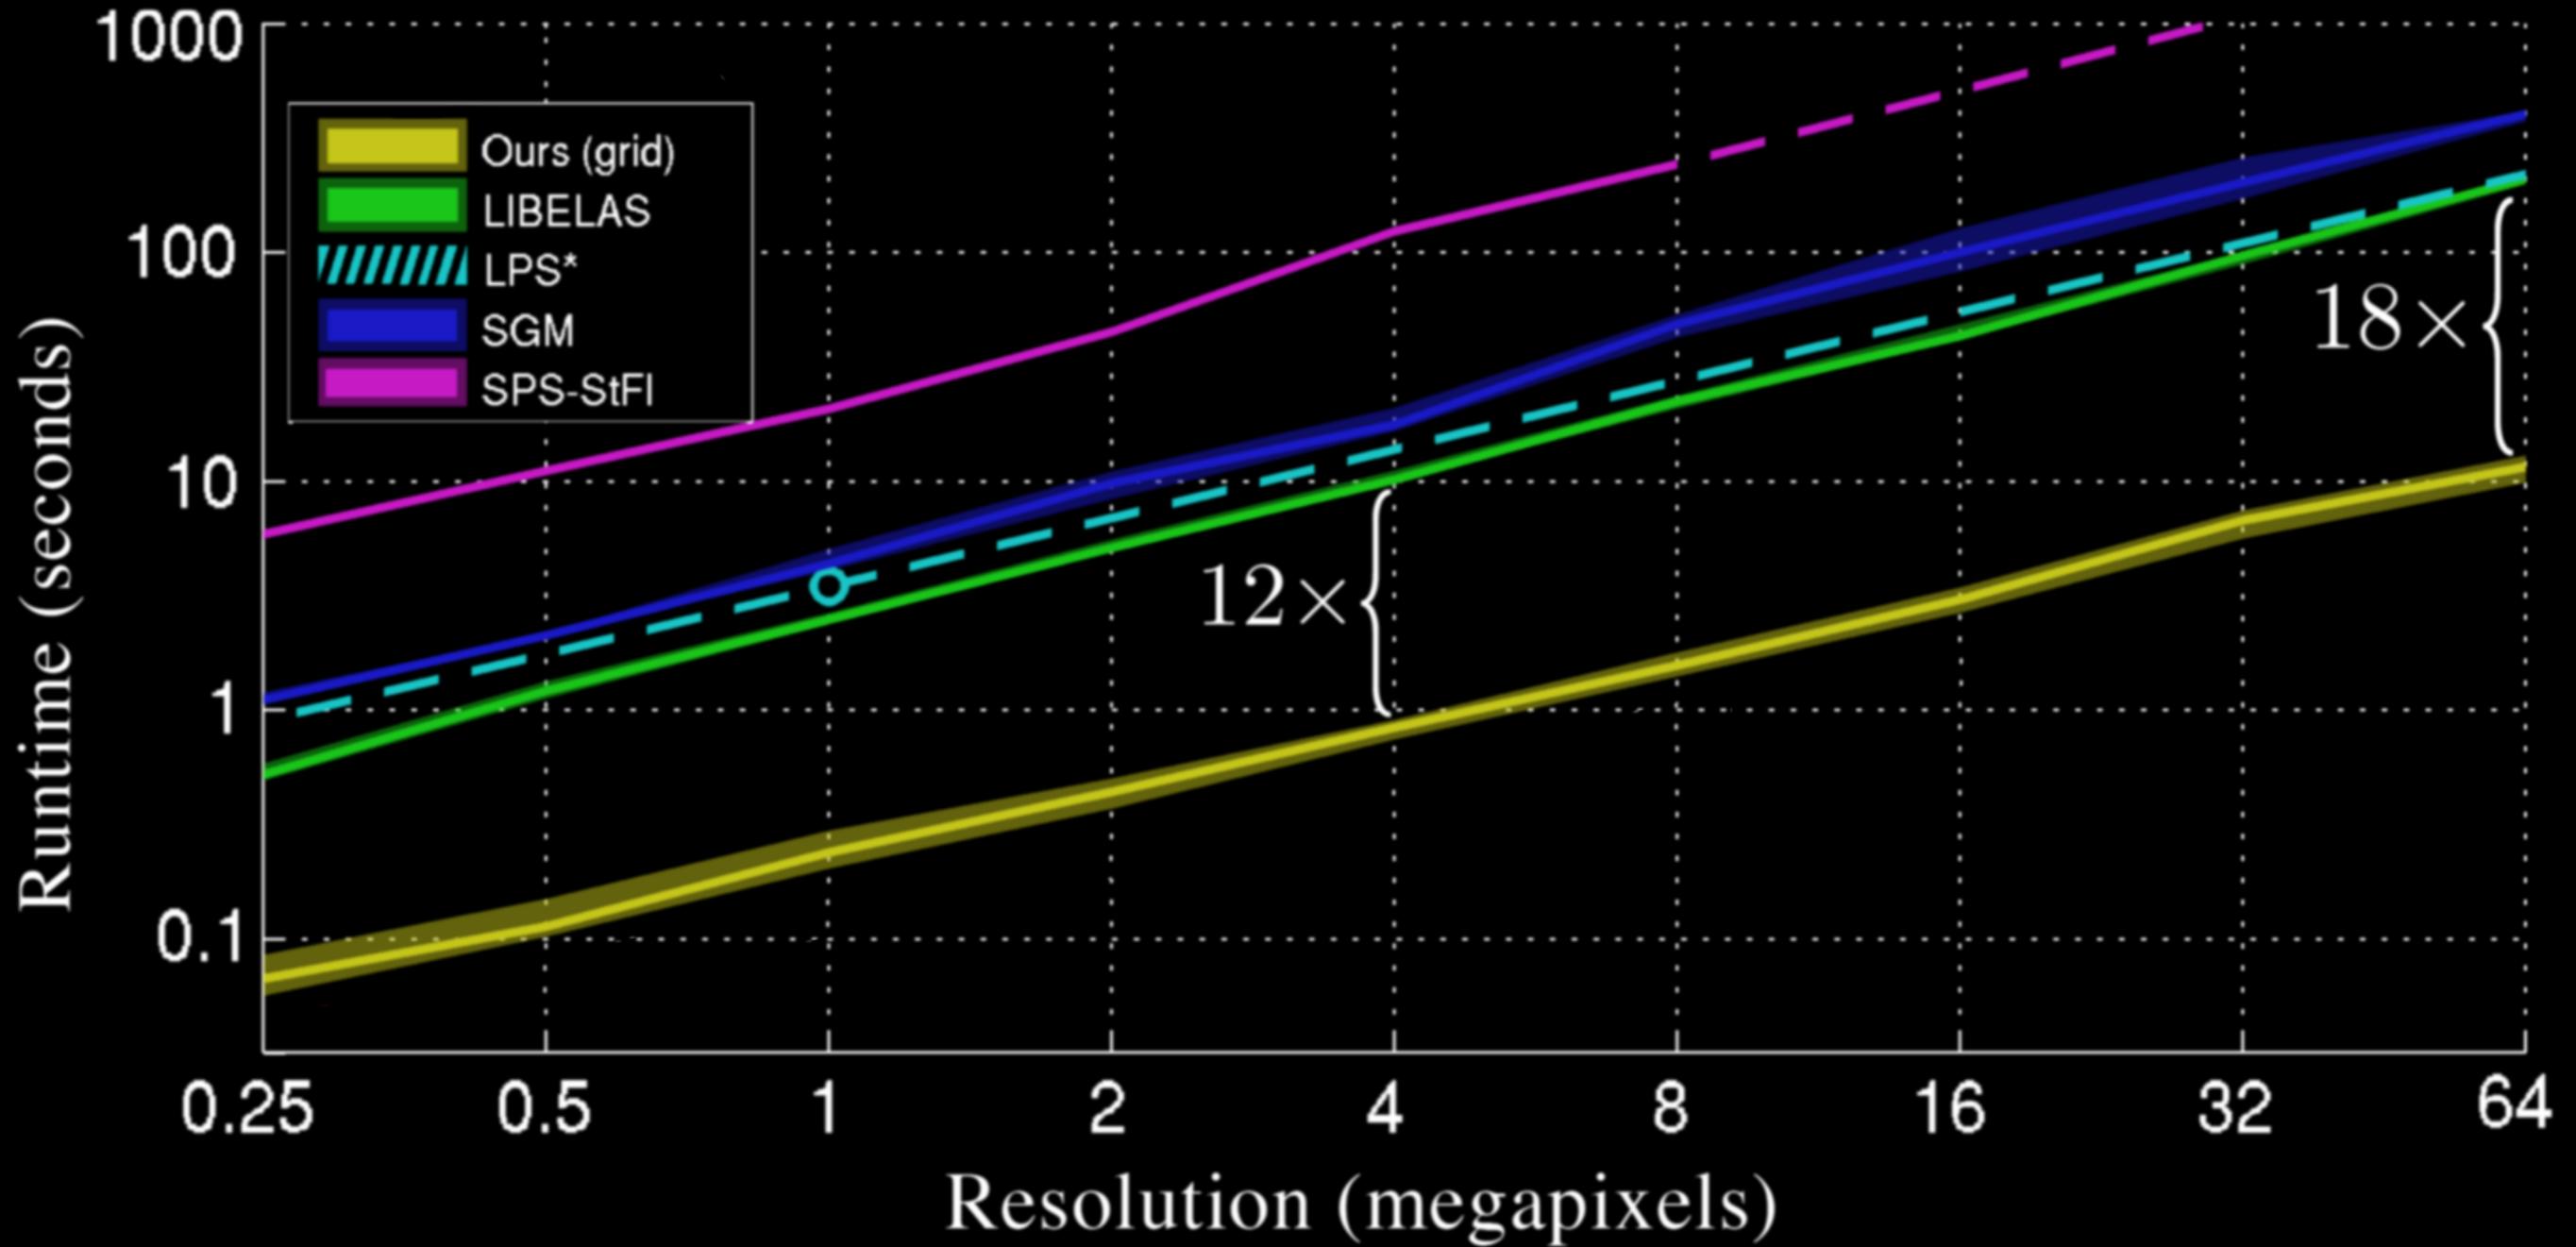
Ours



# Speed



# Speed

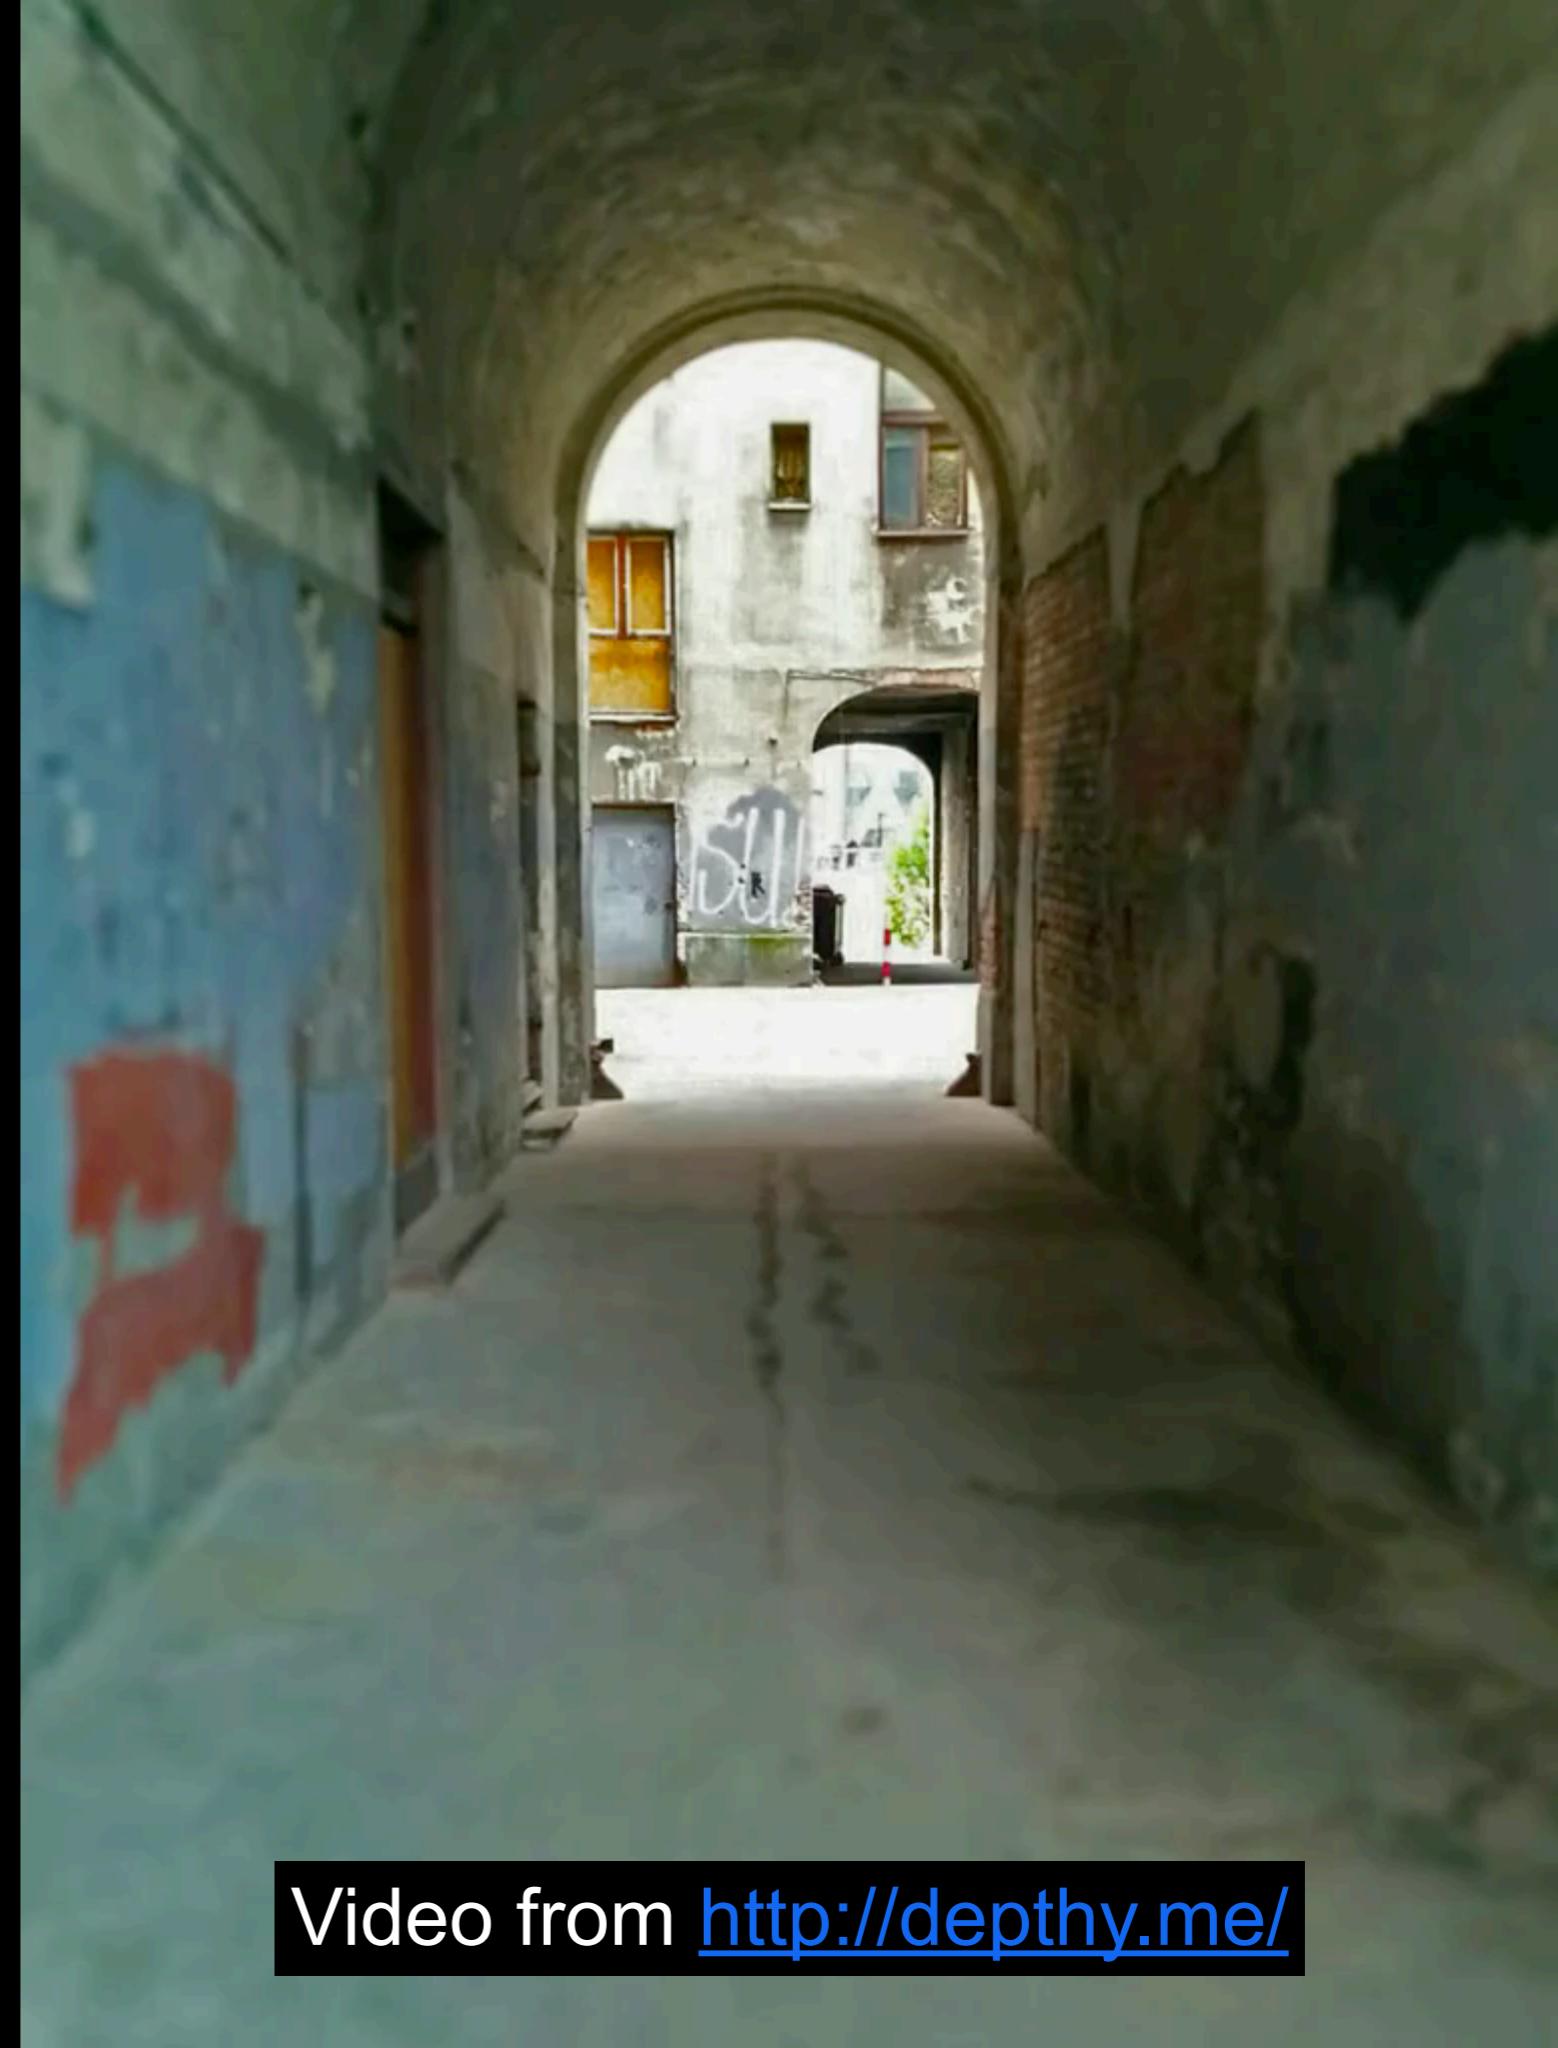


# Lens Blur in the Google Camera App



# Use it!

Depths are embedded in  
the Lens Blur metadata



Video from <http://depty.me/>

# Conclusions

A new way to solve optimization problems based on bilateral filtering

- Fast
- Scalable
- Applicable to many vision / graphics problems

A fast and edge-aware stereo algorithm

- Great for graphics / mobile applications
- Already on millions of phones

# Thanks!

barron@google.com